

---

# Multi-Layer Attention is the Amplifier of Demonstration Effectiveness

---

**Dingzirui Wang<sup>1</sup> Xuanliang Zhang<sup>1</sup> Keyan Xu<sup>1</sup> Qingfu Zhu<sup>1</sup>**  
**Wanxiang Che<sup>1</sup> Yang Deng<sup>2</sup>**

<sup>1</sup>Harbin Institute of Technology <sup>2</sup>Singapore Management University  
{dzrwang, xuanliangzhang, kyxu, qfzhu, car}@ir.hit.edu.cn  
ydeng@smu.edu.sg

## Abstract

Numerous studies have investigated the underlying mechanisms of in-context learning (ICL) effectiveness to inspire the design of related methods. However, existing work predominantly assumes the effectiveness of the demonstrations provided within ICL, while many research indicates that not all demonstrations are effective, failing to yielding any performance improvement during ICL. Therefore, in this paper, we investigate the reasons behind demonstration ineffectiveness. Our analysis is based on gradient flow and linear self-attention models. By setting the gradient flow to zero, we deduce that a demonstration becomes ineffective if its information has either been learned by the model or is irrelevant to the user query. Furthermore, we demonstrate that in multi-layer models, the disparity in effectiveness among demonstrations is amplified with layer increasing, causing the model to focus more on effective ones. Considering that current demonstration selection methods primarily focus on the relevance to the user query while overlooking the information that the model has already assimilated, we propose a novel method called GRADS, which leverages gradient flow for demonstration selection. We use the magnitude of the gradient flow of the demonstration with respect to a given user query as the criterion, thereby ensuring the effectiveness of the chosen ones. We validate our derivation and GRADS on four prominent LLMs across five mainstream datasets. The experimental results confirm that the disparity in effectiveness among demonstrations is magnified as the model layer increases, substantiating our derivations. Moreover, GRADS achieves a relative improvement of 6.8% on average over the strongest baselines, demonstrating its effectiveness.

## 1 Introduction

In-Context Learning (ICL) is an effective method for enhancing the performance of Large Language Models (LLMs), being widely adapted to various tasks [47, 12]. By providing demonstrations relevant to the user query in the input, it guides the reasoning of LLMs, thereby improving inference performance. Recent years have witnessed many efforts on investigating the internal mechanisms of ICL to explain how LLMs acquire this ability for guiding the design of ICL methods [71]. For example, recent works [67, 32, 45] discuss the convergence and convergence speed of ICL, while some other works [36, 22, 8] study the function of each module of Transformer [48] during ICL.

As the performance of LLMs continues to improve, a phenomenon emerges that there exist ineffective demonstrations, which cannot enhance reasoning performance during ICL [11, 52]. However, existing research on the mechanisms of ICL is predominantly based on the assumption that the given demonstrations are effective, which limits the exploration of how to enhance the performance of ICL. Therefore, in this paper, we aim to answer the following research questions (RQs): *1) What makes*

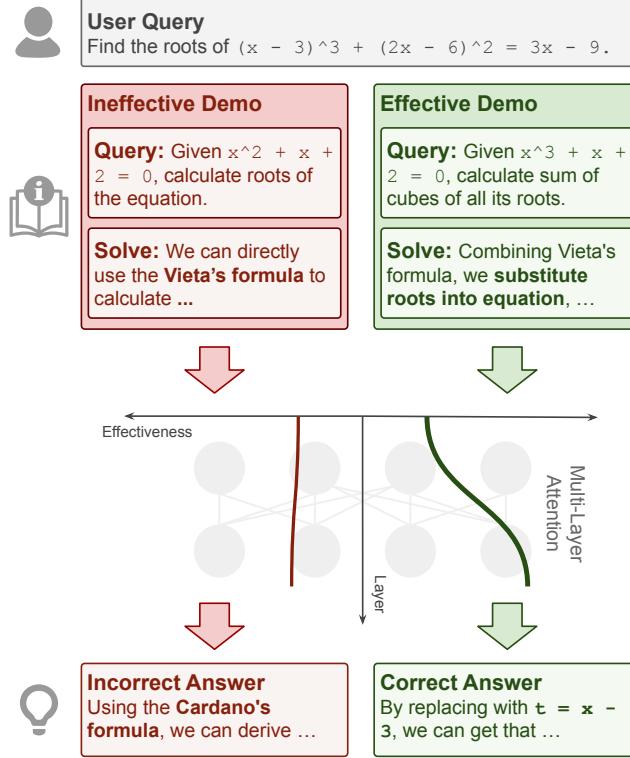


Figure 1: The gradient flow comparison of effective (green) and ineffective demonstrations (red). For the ineffective demonstration, with the increase of model layers, the effectiveness remains at a low level. About the effectiveness, the effectiveness is significantly amplified as the layer increases.

*the demonstration ineffective? 2) How LLMs handle the demonstration effectiveness? and 3) How to select effective demonstrations to enhance ICL performance?*

Following prior works [54, 27], we employ the gradient flow approach for our investigation: the greater the gradient flow from the demonstration to the generated answer, the more effective the demonstration is considered. We first study the factors that determine the magnitude of the gradient flow in a single-layer linear self-attention model, thereby providing the reasons for an ineffective demonstration: the information in the demonstration has already been learned by LLMs, or the demonstration is irrelevant to the user query (RQ1). We then prove that in a multi-layer linear self-attention model, **the difference in effectiveness among demonstrations is amplified along the increase of the number of layers**, resulting in the ignorance of information from less effective demonstrations (RQ2), as shown in Figure 1. Specifically, if a demonstration is more effective than the other, the ratio of their corresponding gradient flows increases with the number of layers, meaning the disparity in their magnitudes grows.

Considering that existing demonstration selection methods mainly focus on the relevance between the demonstration and the query, which could select ineffective demonstrations [59, 5, 51]. Therefore, based on the above discussion, we propose GRADS, which selects the demonstration that maximizes the gradient flow with the given query (RQ3). Measure whether LLMs use the information in a given demonstration through the gradient flow, so as to ensure that the demonstrations provided are effective in the reasoning process, thus enhancing the performance of ICL.

To validate our conclusions and the effectiveness of GRADS, we conduct experiments on five mainstream datasets and four mainstream models. Experimental analysis shows that during ICL, the gradient flow ratio between effective and ineffective demonstrations increases with the number of model layers, proving that the multi-layer Transformer indeed acts as an amplifier of demonstration effectiveness. Furthermore, the experimental results of GRADS indicate that compared to the best demonstration selection baseline, GRADS achieves a relative improvement of 6.8% on average, proving its effectiveness.

In summary, our contributions include:

- We argue that a demonstration in ICL is ineffective when its information is already learned by the LLMs or is irrelevant to the given user query.
- We theoretically and empirically analyze the internal mechanism of LLMs that the multi-layer structure amplifies the ICL effectiveness between the effective demonstrations and the ineffective ones.
- We propose GRADS, a gradient-flow-based demonstration selection method that ensures the selection of highly effective demonstrations.

## 2 Analysis

RQ	Finding	Evidence
<b>RQ1:</b> What makes the demonstration ineffective?	A demonstration is ineffective if the information it contains has already been learned by LLMs or is irrelevant to the user query.	Equation 2
<b>RQ2:</b> How LLMs handle the demonstration effectiveness?	With deeper layers, the gradient flow disparity between effective and ineffective demonstrations widens, prioritizing to learn from the effective one.	Theorem 1

Table 1: The main research question (RQ), findings, and corresponding evidence of our analysis.

In this section, we discuss why and how LLMs ignore ineffective demonstrations in ICL from a gradient flow perspective. We first provide definitions for necessary mathematical notations and concepts. Then, we discuss the gradient flow in a single-layer linear self-attention network and why LLMs distinguish ineffective demonstrations. Subsequently, we discuss the gradient flow in a multi-layer linear self-attention network and how LLMs ignore ineffective demonstrations. The main findings of our analysis are summarized in Table 1. All proofs in this section are provided in Appendix A.1, and the calculation of the gradient flow follows [54].

### 2.1 Preliminary

In this paper, we primarily focus on the 1-shot setting. Following [67], we denote a demonstration as  $d = (d_x \ d_y)$ , where  $d_x, d_y \in \mathbb{R}^e$  represent the input and output embedding vector of the demonstration, respectively. Here,  $e$  is the embedding dimension. We denote a user query as  $q = (q_x \ q_y)$ , where  $q_x \in \mathbb{R}^e$  is the query input embedding vector, and we set  $q_y = 0$  indicates that the corresponding answer to the query is not provided in the input. We denote the complete network input as  $E = (d \ q)$ . In this paper, we use  $\|M\|$  to denote the Frobenius norm [62] of a given matrix  $M$ .

Following the previous work [67], we define the linear self-attention network (LSA) as:

$$f_{LSA}(E; \theta) = E + W^{PV} E \cdot \frac{E^T W^{KQ} E}{\rho} \quad (1)$$

$W^{PV} \in \mathbb{R}^{e \times e}$  is the combined projection and value matrix of Attention, and  $W^{KQ} \in \mathbb{R}^{e \times e}$  is the combined key and query matrix, where the specific definitions are consistent with [67].  $\rho$  is the normalization coefficient, where we set  $\rho = 1$  in this paper following [67]. We denote all network parameters as  $\theta = \{W^{PV}, W^{KQ}, \rho\}$ . It can be observed that Equation 1 is the result of replacing the activation function in a single-layer attention network of a Transformer with a linear function. Specifically, we denote  $\hat{q}_y(E; \theta)$  as the predicted answer of the given  $E$  and  $\theta$ , abbreviated as  $\hat{q}_y$ , which is the vector corresponding to the last column of the output from  $f_{LSA}(E; \theta)$ .

Following previous work [54], for a multivariate function  $f$  and one of its independent variables  $x$ , we define the magnitude of the gradient flow of  $f$  with respect to  $x$  as the partial derivative of  $f$  with respect to  $x$ , which is  $\nabla_x f = \frac{\partial f}{\partial x}$ . The magnitude of the gradient flow measures the influence of the change in  $x$  on  $f$ . Specifically, in the context of ICL, the gradient flow of the output  $\hat{q}_y$  with respect to an input demonstration  $d$  reflects the contribution of that demonstration to the answer. Consequently, it can indicate how much information from the demonstration is utilized in generating the answer.

## 2.2 Single-Layer Linear Self-Attention Network

We first analyze the gradient flow in a single-layer network. It can be shown that the contribution of the input demonstration to the gradient flow for answer generation is:

$$\nabla_d \hat{q}_y = (W^{PV} d)_y (q^\top W^{KQ})^\top + (d^\top W^{KQ} q) W_y^{PV} \quad (2)$$

Equation 2 formally defines the gradient flow value of the single layer LSA, which presents how much information from the demonstration is used for answer generation. We can see that, given a query  $q$ , the factors determining the magnitude of the equation can be divided into two categories: (i) The similarity between the demonstration and the user query ( $d^\top W^{KQ} q$ ). (ii) Whether the model has already learned the information in the demonstration ( $W^{PV} d$ ). Therefore, the determination of whether to use the information from a given demonstration mainly depends on the similarity between the demonstration and the query, and whether the information in the demonstration has already been learned by the model.

Based on the discussion above, as the base of the following discussion, we propose to use the parameters in Equation 2, formally define the effectiveness of demonstrations with the given query and model as follows:

**Definition 1** (Demonstration Effectiveness). *Given a query  $q$  and model parameters  $\theta$ , if two demonstrations  $d_1$  and  $d_2$  satisfy that:*

$$\begin{aligned} \|W^{PV} d_1\| &\geq \|W^{PV} d_2\| \\ \|d_1^\top W^{KQ} q\| &\geq \|d_2^\top W^{KQ} q\| \end{aligned}$$

*then we say that  $d_1$  is more effective than  $d_2$  with respect to  $q$  and  $\theta$ , denoted as:*

$$d_1 \succcurlyeq_{q; \theta} d_2$$

In Definition 1, we require a demonstration to be more effective than the other if it contains more knowledge that the model has not learned, and it is more similar to the query. If only one of these conditions holds, it is difficult to compare the effectiveness of the demonstration because it is not possible to determine whether the knowledge or the similarity to the query has a greater impact on performance. We experimentally evaluate the impact of two factors on ICL performance in Figure 2.

## 2.3 Multi-Layer Linear Self-Attention Network

Next, we discuss the gradient flow of the multi-layer linear self-attention network. Let  $L$  be the total number of layers in the network. We denote the input to the  $l$ -th layer as  $E^{(l)} = (d^{(l)} \quad q^{(l)})$  and its parameters as  $\theta^{(l)}$ , the corresponding predicted answer of  $l$ -th layer is  $\hat{q}^{(l)}$ . Since in the multi-layer LSA, the output of  $(l-1)$ -th layer is the input of the  $l$ -th layer:

$$E^{(l)} = f_{LSA}(E^{(l-1)}; \theta^{(l-1)}) \quad (3)$$

According to the chain rule, we can derive that:

$$\frac{\partial \hat{q}_y^{(L)}}{\partial d^{(0)}} = \frac{\partial \hat{q}_y^{(L)}}{\partial E^{(L-1)}} \times \frac{\partial E^{(L-1)}}{\partial E^{(L-2)}} \times \cdots \times \frac{\partial E^{(1)}}{\partial d^{(0)}} \quad (4)$$

That is, the gradient flow of the whole model is the product of the gradient flow of each layer.

Due to the complexity of the network structure, we only provide a qualitative analysis for the multi-layer model. Based on Equation 2, we know that the magnitude of each term in the chain rule is positively correlated with the demonstration effectiveness, i.e.:

**Lemma 1.** *Let an  $L$ -layer linear self-attention (LSA) network be given. For every layer  $l = 1, \dots, L$ , assume there exist strictly increasing functions*

$$g_l : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}, \quad h_l : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0},$$

*such that for every demonstration  $d$  and query  $q$ , the following equalities hold:*

$$\|W^{PV, (l)} d^{(l)}\| = g_l(\|W^{PV, (l-1)} d^{(l-1)}\|) \quad (5)$$

$$\|(d^{(l)})^\top W^{KQ, (l)} q^{(l)}\| = h_l(\|(d^{(l-1)})^\top W^{KQ, (l-1)} q^{(l-1)}\|) \quad (6)$$

*If two demonstrations satisfy  $d_1 \succcurlyeq_{q; \theta^{(0)}} d_2$  at the input layer, then for every  $l = 1, \dots, L$ , we have  $d_1^{(l)} \succcurlyeq_{q; \theta^{(l)}} d_2^{(l)}$ .*

The condition in Lemma 1 assumes that each layer of the model has a consistent effect on its input. Based on Lemma 1, we know that for each layer, a more effective input demonstration results in a larger corresponding gradient flow on all layers. It is worth noting that for the same demonstration  $d$ , Lemma 1 does not guarantee that the gradient flow at layer  $l$  is always greater than that at layer  $l - 1$ . This is because the magnitude of the gradient flow also depends on the parameters  $\theta^{(l)}$  of each layer, so it cannot be guaranteed that the gradient flow is monotonically increasing across each layer.

Based on Lemma 1, we can deduce that since the gradient flow is a partial derivative, a larger gradient flow leads to a greater change in the output of subsequent layers. Therefore, as the number of model layers increases, the difference in gradient flow between effective and ineffective demonstrations also becomes larger.

**Theorem 1** (Multi-Layer Attention is the Amplifier of Demonstration Effectiveness). *For a given user query  $q$  and a model  $\theta$ , let  $d_1$  and  $d_2$  be two demonstrations with corresponding inputs  $E_1$  and  $E_2$ . If the condition of Lemma 1 holds, for any  $L \geq l_1 > l_2 \geq 1$ , we can draw that the following inequalities hold:*

$$\frac{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_1)}(E_1; \theta)\|}{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_1)}(E_2; \theta)\|} \geq \frac{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_2)}(E_1; \theta)\|}{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_2)}(E_2; \theta)\|}$$

Theorem 1 shows that as the number of layers increases, the gap between the cumulative gradient flows of different demonstrations also becomes increasingly large. This indicates that the multi-layer LSA acts as an amplifier for demonstration effectiveness.

### 3 Methodology

In this section, we introduce GRADS, our demonstration selection method based on gradient flow, selecting the demonstration with the largest flow to the given query as the selection result. The prompt we used is shown in Appendix A.3. Considering Theorem 1, in actual selection, we use the gradient flow of the last layer as the selection metric. This is because the differences between demonstrations are sufficiently amplified, allowing for a better distinction between effective and ineffective demonstrations.

However, performing a full inference pass on the model to obtain the gradient flow for each user query results in low computational efficiency. From Equation 2, it can be observed that in the calculation of the gradient flow, the computations for the demonstration and the user query are relatively independent. Therefore, we can first compute the encoded vectors for the demonstrations and the user query separately, and then use these results to calculate the magnitude of the gradient flow through matrix operations. This approach significantly reduces the computational overhead, thereby enhancing the efficiency of demonstration selection.

Specifically, for a given demonstration pool, we first input each demonstration individually to extract the encoding result of the final layer as  $\hat{d}$ . Then, for each user query, we also input it to extract the encoding result of the final layer as  $\hat{q}$ . Subsequently, the computed  $\hat{d}$  and  $\hat{q}$  are substituted into Equation 2 to obtain  $\partial_d \hat{q}_y^{(L)}$ . Since  $\partial_d \hat{q}_y^{(L)}$  is a  $2 \times e$  matrix, containing the gradient flows for both the input and output parts of the demonstration, and considering that both parts are equally important, we use  $\|\partial_d \hat{q}_y^{(L)}\|$  as the metric for demonstration selection to balance their importance.

**Efficiency of GRADS** The computational cost of GRADS is mainly divided into two parts: the offline computation of the encoding result for each demonstration, and the online retrieval of the demonstration and subsequent inference for a query. Let  $D = \{d_1, \dots, d_n\}$  represent the entire demonstration pool, and let  $\mathcal{M}_\theta(x)$  denote the computational cost of the model  $\mathcal{M}_\theta$  for a given input  $x$ . For a query  $q$ , let the retrieved demonstration be  $d_q$ .

The time complexity of the offline processing for GRADS is:  $O(\sum_{i=1}^n \mathcal{M}_\theta(d_i))$ . Although the offline processing cost is relatively high, the pre-computation of demonstration encodings is done offline and thus does not affect the online user query process.

The time complexity for online processing is:

$$O(\mathcal{M}_\theta(q) + 4e^2 + \mathcal{M}_\theta(q + d_q)) \quad (7)$$

In Equation 7, the first term represents encoding the user query, the second term corresponds to calculating Equation 2, and the third term is for generating the answer to the user query based on the retrieved demonstration. Considering that  $\mathcal{M}_\theta$  is positively correlated with the input length and the complexity of calculating Equation 2 is significantly lower than that of a full model inference  $\mathcal{M}_\theta$ , the online processing time complexity simplifies to:

$$O(\mathcal{M}_\theta(q + d_q)) \quad (8)$$

This is equivalent to the time complexity of direct 1-shot inference, which demonstrates the high efficiency of GRADS.

## 4 Experiment

Our experiments are primarily divided into three parts: (i) Introduction of the experimental settings and baselines. (ii) Verification of Theorem 1 and its related corollaries. (iii) Validation of effectiveness and the impact of different factors on GRADS.

### 4.1 Experiment Setup

**Dataset** To thoroughly validate our analytic conclusions and the effectiveness of our proposed method, we conduct experiments on five mainstream datasets that span various tasks and domains, including: (i) math: GSM8K [10] and MATH [17]; (ii) reasoning: ARC-Challenge [63] and MMU-Pro [58]; (iii) sentiment analysis: Amazon Review [34]. Detailed descriptions of these datasets are provided in Appendix A.4. We employ Exact Match (EM) as the evaluation metric across all datasets.

**Model** We validate our discovery and method on four LLMs: Llama2-7b [20], Llama3.1-8b [1], Deepseek-R1-Distilled-Llama3.1-8b (Llama-R1-8b) [11], and Qwen3-8b [2]<sup>1</sup>. Our selection of models encompasses various scales, series, and capabilities for generating long chains of thought (Long-CoT) [6]. This diverse set allows for a robust verification of our conclusions and enables a comparative study on the influence of model characteristics on our findings.

**Baseline** We compare our method against three classes of demonstration synthesis baselines following [51]: (i) n-gram-based methods (BM25) [41, 23], (ii) vector similarity-based methods (Cosine) [65], and (iii) LLM-based methods (MMR [66], MoD [56]). Detailed introductions and configurations for each baseline are provided in Appendix A.5.

**Implementation Detail** For the mechanism analysis experiments, we employ a 1-shot setting to align with the setup in our theoretical analysis. For the validation of GRADS, we follow previous works [51] and adopt a 3-shot setting to ensure comparability of the final performance. Following [11], we set the maximum generation length to 32768, and for each question, we sample a single answer. Our experiments are performed on a single A100-80G GPU, with the selection and inference for each dataset taking approximately one hour on average.

### 4.2 Experiment of Mechanism Analysis

#### 4.2.1 The Factors Making Demonstrations Ineffective

To verify the condition for low gradient flow in the discussion regarding Equation 2, we record the statistics of the model performance as a function of the changes in the demonstration relevance ( $d^\top W^{KQ} q$ ) and learned knowledge ( $W^{PV} d$ ). The experimental results are shown in Figure 2. From the figure, we can observe that: (i) As the relevance of the demonstrations and the knowledge learned from them increase, the number of correctly solved data points also increases, which verifies the conclusion from Equation 2 that the performance of ICL is positively correlated with these two factors. (ii) ICL only begins to correctly solve the given user queries after the relevance of the demonstrations and the knowledge learned from them surpass a certain threshold, which indicates that a sufficient amount of effective information that is relevant to the user query is necessary to enable the model to perform correct reasoning.

---

<sup>1</sup>We utilize the thinking mode of Qwen3 for a comprehensive evaluation.

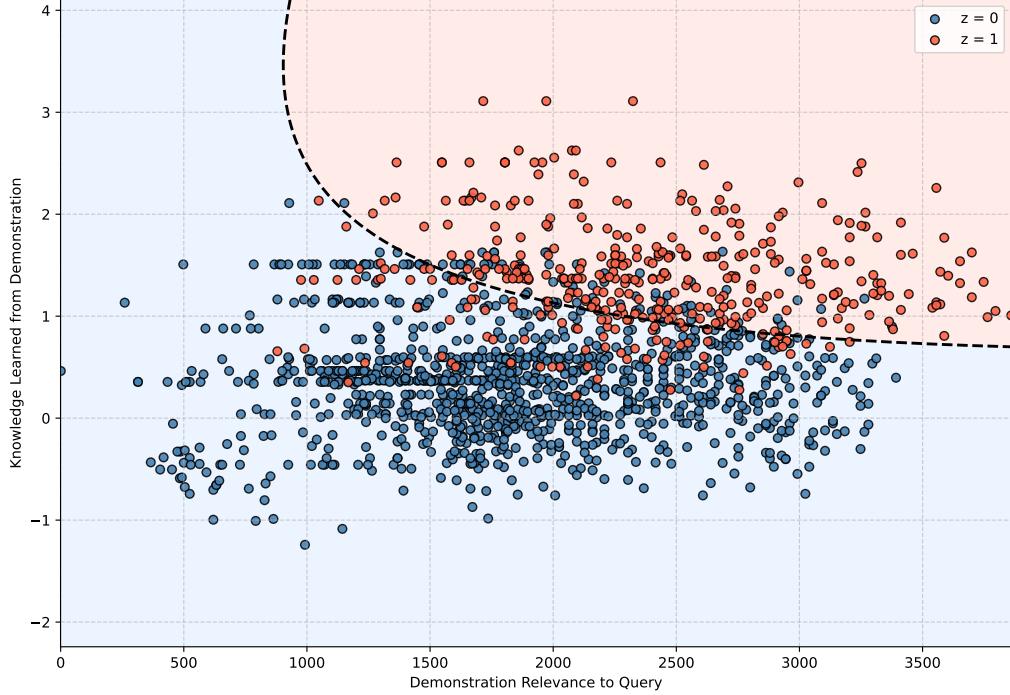


Figure 2: The ICL performance of Llama3.1-8b across all datasets with different demonstration relevance (X-axis) and unlearned knowledge within the demonstration (Y-axis). Red points denote correct prediction and blue points incorrect predictions. The dashed line represents the decision boundary generated with polynomial logistic regression [4].

#### 4.2.2 Gradient Flow is Amplified as the Layer Increasing

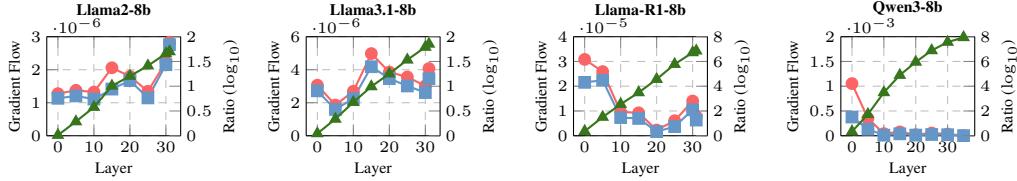


Figure 3: The average gradient flow (left y-axis) and the ratio  $\frac{\|\nabla_{d(0)} \hat{q}_y^{(l1)}(E_1; \theta)\|}{\|\nabla_{d(0)} \hat{q}_y^{(l1)}(E_2; \theta)\|}$  of Theorem 1 (right y-axis) cross different datasets under the  $i$ -th layer of each model, where  $E_1$  denotes the input of the effective demonstrations and  $E_2$  denotes the ineffective ones. Red points denote the average gradient flow of the effective demonstrations, blue points denote the ineffective demonstrations, and green points denote the ratio.

To validate the correctness of Lemma 1 and Theorem 1, we compute the gradient flow on different settings. For a given dataset, we first select the data points that are incorrectly predicted under the 0-shot setting. From this subset, we then separate the data into two groups based on the 1-shot performance: those that become correctly predicted (*effective demonstrations*) and those that remain incorrectly predicted (*ineffective demonstrations*). We then compute the average gradient flow for each layer across these two groups. The number of selected effective and ineffective data points is detailed in Appendix A.6.1.

The experimental results are presented in Figure 3. From the figure, we can observe the following: (i) Across all settings, the average gradient flow of effective demonstrations is stronger than that of ineffective demonstrations at every layer, which validates the conclusion of Lemma 1. Furthermore, the ratio of gradient flow between effective and ineffective demonstrations increases with the layer depth, which supports Theorem 1. (ii) The magnitude of the gradient flow, as well as the ratio, is

significantly more pronounced in the Long-CoT model compared to other models. This suggests that such models are more adept at capturing information from demonstrations and can better distinguish between effective and ineffective demonstrations, showing a higher sensitivity to the useful information contained within the demonstrations. (iii) In all models, there are specific layers where the gradient flow exhibits a significant increase. This indicates that the learning from demonstrations is primarily concentrated in these particular layers. While the specific layers differ across models, a strong gradient flow is consistently observed in the final few layers, suggesting that the model pays special attention to the information in the demonstrations when generating the output.

### 4.3 Experiment of GRADS

#### 4.3.1 GRADS is More Effective than Baselines

Model	Method	GSM8K	MATH	ARC-C	MMLU-Pro	Amazon
Llama2-7b	Zero	12.7	5.0	34.6	14.5	28.5
	BM25	25.7	8.4	44.8	17.1	32.0
	Cosine	25.1	7.2	45.1	17.5	37.0
	MMR	24.9	8.0	45.7	17.7	33.5
	MoD	25.1	7.2	45.1	17.5	37.0
	GRADS	<b>26.7</b>	<b>9.6</b>	<b>46.5</b>	<b>18.8</b>	<b>37.5</b>
Llama3.1-8b	Zero	83.7	47.0	82.0	50.4	63.5
	BM25	84.5	44.2	84.6	52.7	69.5
	Cosine	85.0	47.0	85.2	52.3	69.0
	MMR	84.2	45.0	85.3	52.0	66.5
	MoD	85.0	47.0	85.2	51.9	69.0
	GRADS	<b>85.6</b>	<b>48.2</b>	<b>86.3</b>	<b>56.0</b>	<b>70.0</b>
Llama-R1-8b	Zero	86.1	75.4	83.5	58.2	61.5
	BM25	80.1	74.2	84.9	52.7	65.0
	Cosine	86.0	74.6	84.9	55.9	65.5
	MMR	85.6	72.8	84.4	59.0	66.0
	MoD	86.0	74.6	84.4	56.9	65.5
	GRADS	<b>87.2</b>	<b>75.6</b>	<b>85.7</b>	<b>59.5</b>	<b>67.0</b>
Qwen3-8b	Zero	93.9	76.2	89.2	64.9	62.5
	BM25	93.2	76.8	92.0	64.0	65.0
	Cosine	93.9	77.4	90.4	65.0	68.5
	MMR	93.1	77.4	90.8	65.5	74.5
	MoD	93.9	78.2	90.9	64.8	69.0
	GRADS	<b>94.2</b>	<b>79.4</b>	<b>92.6</b>	<b>68.5</b>	<b>75.0</b>

Table 2: The performance of GRADS compared with different baselines. ARC-C denotes ARC-Challenge, Amazon denotes Amazon Review. The best result under each setting is marked in **bold**.

To validate the effectiveness of GRADS, we conduct a comparative analysis against several baselines, with the experimental results presented in Table 2. The result shows that GRADS achieves a relative improvement of 6.8% on average over the best-performing baseline cross different setting, which substantiates its efficacy and generalizability. Furthermore, a deeper analysis of the results reveals several key observations:

(i) *From a methodological perspective*: In many settings, the performance of MMR and MoD does not exceed that of simpler methods like BM25, and in some cases, is even inferior to the zero-shot approach. This suggests that methods based on the similarity between demonstrations and the user query do not guarantee an enhancement in ICL performance due to that the model could have already been exposed to the information present in the demonstrations, and irrelevant information within these demonstrations could consequently mislead the model reasoning process. In contrast, GRADS ensures that the information in the selected demonstrations is effectively utilized by the model, thereby securing the effectiveness of the selected demonstrations.

(ii) *From a model perspective*: The most significant performance improvement from GRADS is observed on Llama3.1-8b. This is likely because this model does not employ Long-CoT, rendering it less capable of effectively leveraging the useful information within the demonstrations. Consequently, its performance is more dependent on the quality of the demonstrations compared to Llama-R1-8b and Qwen3-8b. Conversely, the performance gain on Llama2-7b is relatively modest. We attribute this to the limited knowledge base of Llama2 [55], which makes the relevance between the demonstrations and the user query a more critical factor for ICL performance. As a result, the performance gap between GRADS and other relevance-based baselines is less pronounced.

(iii) *From a dataset perspective*: GRADS yielded more substantial performance gains on the MATH and MMLU-Pro datasets. This is because these two datasets, compared to others, demand a higher

Model	Dataset	Layer							
		0	5	10	15	20	25	30	31
Llama3.1-8b	GSM8K	83.2	84.4	84.9	85.1	85.2	85.6	85.2	<b>85.6</b>
	MATH	44.4	44.8	45.8	45.2	46.2	47.6	46.8	<b>48.2</b>
	ARC-Challenge	82.1	82.3	82.1	82.8	83.4	84.0	85.1	<b>86.3</b>
	MMLU-Pro	50.8	51.5	53.6	54.8	54.8	55.6	<b>56.2</b>	56.0

Table 3: The performance of GRADS using the gradient flow of different layer, where 31 is the last layer of the model. The best result under each setting is marked in **bold**.

degree of specialized knowledge. Therefore, they are more reliant on the demonstrations to provide knowledge that the model lacks. GRADS, being more effective at selecting demonstrations that contain the requisite knowledge for the model, achieves a more significant performance improvement in these contexts.

#### 4.3.2 The Performance of GRADS is Positively Correlated to Model Layer

To validate Theorem 1 that the disparity in gradient flow between effective and ineffective demonstrations increases with network layer, thereby enhancing the demonstration selection capability, we conduct the experiment on the performance of GRADS using gradient flows from different layers for demonstration selection. The experimental results are presented in Table 3. From the table, we can observe the following: *(i)* Across all datasets, there is a general upward trend in model performance as the number of layers increases, which confirms the conclusion of Theorem 1 that as network layer increases, the difference in effectiveness among demonstrations is amplified, which enables a better selection of effective demonstrations. *(ii)* However, the performance of GRADS does not monotonically increase with the number of layers since the models used in our experiments are more complex than those in our theoretical analysis. For instance, the presence of residual streams [15] can diminish the amplifying effect of multi-layer transformers on the gradient flow, which could lead the model to erroneously select ineffective demonstrations, resulting in a decline in performance.

## 5 Related Works

### 5.1 In-Context Learning

ICL is an effective method for enhancing the performance of LLMs by providing demonstrations related to the user query in the input to improve reasoning performance [47, 12]. Existing ICL research can be categorized into three areas: how to acquire demonstrations, how to select demonstrations, and how to utilize demonstrations. To reduce the cost of manually labeling demonstrations, many methods have been proposed to synthesize demonstrations using LLMs [28], relying on resources such as existing demonstrations [46, 16], related information [7, 51, 9], and data from similar tasks [50] to generate demonstrations relevant to the target task. Demonstration selection primarily focuses on how to select demonstrations from a pool that are relevant to the target query [38, 57]. Early work mainly relies on gram-based methods [23], while recent work leverages model-enhanced processing of information within demonstrations and queries [65, 66, 56]. Demonstration utilization focuses on how to make better use of the selected demonstrations, for example, by adjusting the order of the demonstrations [30, 39] or encoding them into vectors and directly injecting them into the model [21, 53, 24], thereby further enhancing ICL performance while reducing its inference overhead.

However, existing demonstration selection methods primarily focus on the relevance between the demonstration and the query, overlooking the phenomenon that the demonstrations could be ineffective due to that the information in the demonstrations has already been learned by the model. Therefore, we propose GRADS, which is based on gradient flow, to ensure that the selected demonstrations provide a sufficiently large information contribution during inference, thereby enhancing the performance of ICL.

## 5.2 Mechanism of In-Context Learning

Understanding the internal mechanism of ICL can help us better enhance its performance and comprehend the model reasoning process [71]. Existing work on the mechanism of ICL can be divided into four categories: theoretical derivation, model architecture, training data, and inference analysis. Theoretical derivation work primarily focuses on using mathematical proofs to demonstrate the effectiveness of ICL, such as its convergence [61, 19, 64] and convergence rate [45, 13, 18, 49], among which many works adapt the research based on LSA [67, 32, 31]. The study of model architecture mainly discusses the roles of different modules within Transformer architecture in ICL. For example, the Attention mechanism plays a major role in ICL [36, 8, 35], while the MLP layers primarily serve an auxiliary function [22, 33]. Training data analysis examines how the ICL ability emerges during the training process. Current work suggests that the ICL ability mainly originates from the diversity of training tasks [40, 60, 68, 14, 69], and during training, the model gradually transitions from in-weight learning to in-context learning [44, 37]. Inference analysis relies on observing phenomena during ICL inference, including the influence of different factors on ICL and changes in internal computational information [3, 42, 26, 29, 43, 70].

However, existing research on the ICL mechanism assumes that the given demonstrations are effective. In actual inference, there is the phenomenon that demonstrations are ineffective, leading to no performance improvement [11, 52]. Therefore, in this paper, we investigate the ineffective demonstrations, provide reasons for it including the demonstration being irrelevant or its information having already been learned by the model, and propose that multi-layer transformers act as amplifiers of demonstration effectiveness.

## 6 Conclusion

In this paper, we primarily discuss the phenomenon of ineffective demonstrations in ICL. Based on LSA, we first discuss that a demonstration is ineffective because the information it contains has already been learned by the model or is irrelevant to the user query. We then demonstrate that in multi-layer LSA, as the number of model layers increases, the distinction in effectiveness among demonstrations is amplified, leading the model to focus more on the information within relevant demonstrations. To select effective demonstrations for ICL, based on the above discussion, we propose GRADS, which uses gradient flow as a metric to select demonstrations. This ensures the effectiveness of the selected results, thereby enhancing ICL performance. To validate these conclusions, we conduct experiments on four mainstream LLMs and five mainstream datasets, covering various tasks and domains. First, under all experimental settings, analytical experiments show that as the number of model layers increases, the difference in effectiveness among demonstrations is progressively magnified, which corroborates our theoretical derivations. Second, GRADS achieves an average relative performance improvement of 6.8% compared to existing demonstration selection methods, proving the effectiveness of our method.

## References

- [1] et al Aaron Grattafiori, Abhimanyu Dubey. The llama 3 herd of models, 2024.
- [2] et al An Yang, Anfeng Li. Qwen3 technical report, 2025.
- [3] Eric J. Bigelow, Ekdeep S. Lubana, Robert P. Dick, Hidenori Tanaka, and Tomer D. Ullman. In-context learning dynamics with random binary sequences. In *International Conference on Learning Representations (ICLR)*, 2024. OpenReview (ICLR 2024).
- [4] Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, Robert Layton, Jake VanderPlas, Arnaud Joly, Brian Holt, and Gaël Varoquaux. API design for machine learning software: experiences from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, pages 108–122, 2013.
- [5] Juhai Chen, Lichang Chen, Chen Zhu, and Tianyi Zhou. How many demonstrations do you need for in-context learning? In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 11149–11159, Singapore, December 2023. Association for Computational Linguistics.
- [6] Qiguang Chen, Libo Qin, Jinshao Liu, Dengyun Peng, Jiannan Guan, Peng Wang, Mengkang Hu, Yuhang Zhou, Te Gao, and Wanxiang Che. Towards reasoning era: A survey of long chain-of-thought for reasoning large language models, 2025.
- [7] Wei-Lin Chen, Cheng-Kuang Wu, Yun-Nung Chen, and Hsin-Hsi Chen. Self-ICL: Zero-shot in-context learning with self-generated demonstrations. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 15651–15662, Singapore, December 2023. Association for Computational Linguistics.
- [8] Xingwu Chen, Lei Zhao, and Difan Zou. How transformers utilize multi-head attention in in-context learning? a case study on sparse linear regression. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [9] Zihan Chen, Song Wang, Zhen Tan, Jundong Li, and Cong Shen. MAPLE: Many-shot adaptive pseudo-labeling for in-context learning. In *Forty-second International Conference on Machine Learning*, 2025.
- [10] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021.
- [11] DeepSeek-AI. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025.
- [12] Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Jingyuan Ma, Rui Li, Heming Xia, Jingjing Xu, Zhiyong Wu, Baobao Chang, Xu Sun, Lei Li, and Zhifang Sui. A survey on in-context learning. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 1107–1128, Miami, Florida, USA, November 2024. Association for Computational Linguistics.
- [13] Deqing Fu, Tianqi Chen, Robin Jia, and Vatsal Sharan. Transformers learn to achieve second-order convergence rates for in-context linear regression. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [14] Chase Goddard, Lindsay M. Smith, Vudtiwat Ngampruetikorn, and David J. Schwab. When can in-context learning generalize out of task distribution? In *International Conference on Machine Learning (ICML)*, 2025. OpenReview (ICML 2025).
- [15] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016.

- [16] Wei He, Shichun Liu, Jun Zhao, Yiwen Ding, Yi Lu, Zhiheng Xi, Tao Gui, Qi Zhang, and Xuanjing Huang. Self-demos: Eliciting out-of-demonstration generalizability in large language models. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3829–3845, Mexico City, Mexico, June 2024. Association for Computational Linguistics.
- [17] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)*, 2021.
- [18] Jianhao Huang, Zixuan Wang, and Jason D. Lee. Transformers learn to implement multi-step gradient descent with chain of thought. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [19] Yu Huang, Yuan Cheng, and Yingbin Liang. In-context convergence of transformers, 2024.
- [20] et al Hugo Touvron, Louis Martin. Llama 2: Open foundation and fine-tuned chat models, 2023.
- [21] Dongfang Li, zhenyu liu, Xinshuo Hu, Zetian Sun, Baotian Hu, and Min Zhang. In-context learning state vector with inner and momentum optimization. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [22] Hongkang Li, Meng Wang, Songtao Lu, Xiaodong Cui, and Pin-Yu Chen. How do nonlinear transformers learn and generalize in in-context learning? In *Forty-first International Conference on Machine Learning*, 2024.
- [23] Xiaonan Li, Kai Lv, Hang Yan, Tianyang Lin, Wei Zhu, Yuan Ni, Guotong Xie, Xiaoling Wang, and Xipeng Qiu. Unified demonstration retriever for in-context learning. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4644–4668, Toronto, Canada, July 2023. Association for Computational Linguistics.
- [24] Zhuwei Li, Zihao Xu, Ligong Han, Yunhe Gao, Song Wen, Di Liu, Hao Wang, and Dimitris N. Metaxas. Implicit in-context learning. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [25] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *The Twelfth International Conference on Learning Representations*, 2024.
- [26] Ziqian Lin and Kangwook Lee. Dual operating modes of in-context learning. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, pages 30135–30188. PMLR, 2024. PMLR (ICML 2024).
- [27] Yiting Liu and Zhi-Hong Deng. Iterative vectors: In-context gradient steering without back-propagation. In *Forty-second International Conference on Machine Learning*, 2025.
- [28] Lin Long, Rui Wang, Ruixuan Xiao, Junbo Zhao, Xiao Ding, Gang Chen, and Haobo Wang. On LLMs-driven synthetic data generation, curation, and evaluation: A survey. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 11065–11082, Bangkok, Thailand, August 2024. Association for Computational Linguistics.
- [29] Quanyu Long, Yin Wu, Wenya Wang, and Sinno Jialin Pan. Does in-context learning really learn? rethinking how large language models respond and solve tasks via in-context learning. In *Conference on Learning and Modeling (COLM)*, 2024. OpenReview (COLM 2024).
- [30] Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. Fantastically ordered prompts and where to find them: Overcoming few-shot prompt order sensitivity. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8086–8098, Dublin, Ireland, May 2022. Association for Computational Linguistics.

- [31] Yue Lu, Mary Letey, Jacob A Zavatone-Veth, Anindita Maiti, and Cengiz Pehlevan. In-context learning by linear attention: Exact asymptotics and experiments. In *NeurIPS 2024 Workshop on Mathematics of Modern Machine Learning*, 2024.
- [32] Arvind V. Mahankali, Tatsunori Hashimoto, and Tengyu Ma. One step of gradient descent is provably the optimal in-context learner with one layer of linear self-attention. In *The Twelfth International Conference on Learning Representations*, 2024.
- [33] Alex Nguyen and Gautam Reddy. Differential learning kinetics govern the transition from memorization to generalization during in-context learning. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [34] Jianmo Ni, Jiacheng Li, and Julian McAuley. Justifying recommendations using distantly-labeled reviews and fine-grained aspects. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 188–197, Hong Kong, China, November 2019. Association for Computational Linguistics.
- [35] Kazusato Oko, Yujin Song, Taiji Suzuki, and Denny Wu. Pretrained transformer efficiently learns low-dimensional target functions in-context. In *NeurIPS*, 2024.
- [36] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. In-context learning and induction heads, 2022.
- [37] Core Francisco Park, Ekdeep Singh Lubana, and Hidenori Tanaka. Competition dynamics shape algorithmic phases of in-context learning. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [38] Keqin Peng, Liang Ding, Yancheng Yuan, Xuebo Liu, Min Zhang, Yuanxin Ouyang, and Dacheng Tao. Revisiting demonstration selection strategies in in-context learning. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9090–9101, Bangkok, Thailand, August 2024. Association for Computational Linguistics.
- [39] Kha Pham, Hung Le, Man Ngo, and Truyen Tran. Rapid selection and ordering of in-context demonstrations via prompt embedding clustering. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [40] Allan Raventós, Mansheej Paul, Feng Chen, and Surya Ganguli. Pretraining task diversity and the emergence of non-bayesian in-context learning for regression. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. OpenReview (NeurIPS 2023).
- [41] Stephen Robertson and Hugo Zaragoza. The probabilistic relevance framework: Bm25 and beyond. *Found. Trends Inf. Retr.*, 3(4):333–389, April 2009.
- [42] Zhenmei Shi, Zhouyan Xu, Junyi Wei, and Yingyu Liang. Why larger language models do in-context learning differently? In *International Conference on Machine Learning (ICML)*, 2024. OpenReview (ICML 2024).
- [43] Suzanna Sia, David Mueller, and Kevin Duh. Where does in-context learning happen in large language models? In *NeurIPS (Poster)*, 2024. OpenReview (NeurIPS 2024).
- [44] Aaditya K Singh, Ted Moskovitz, Sara Dragutinović, Felix Hill, Stephanie C.Y. Chan, and Andrew M Saxe. Strategy competition explains the emergence and transience of in-context learning. In *Forty-second International Conference on Machine Learning*, 2025.
- [45] Matthew Smart, Alberto Bietti, and Anirvan M. Sengupta. In-context denoising with one-layer transformers: Connections between attention and associative memory retrieval. In *Forty-second International Conference on Machine Learning*, 2025.

- [46] Yi Su, Yunpeng Tai, Yixin Ji, Juntao Li, Yan Bowen, and Min Zhang. Demonstration augmentation for zero-shot in-context learning. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 14232–14244, Bangkok, Thailand, August 2024. Association for Computational Linguistics.
- [47] et al Tom B. Brown, Benjamin Mann. Language models are few-shot learners, 2020.
- [48] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In I. Guyon, U. Von Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
- [49] Max Vladymyrov, Johannes von Oswald, Mark Sandler, and Rong Ge. Linear transformers are versatile in-context learners. In *ICML 2024 Workshop on In-Context Learning*, 2024.
- [50] Dingzirui Wang, Xuanliang Zhang, Qiguang Chen, Longxu Dou, Xiao Xu, Rongyu Cao, YINGWEI MA, Qingfu Zhu, Wanxiang Che, Binhu Li, Fei Huang, and Yongbin Li. In-context transfer learning: Demonstration synthesis by transferring similar tasks, 2025.
- [51] Dingzirui Wang, Xuanliang Zhang, Keyan Xu, Qingfu Zhu, Wanxiang Che, and Yang Deng. V-synthesis: Task-agnostic synthesis of consistent and diverse in-context demonstrations from scratch via v-entropy, 2025.
- [52] Dingzirui Wang, Xuanliang Zhang, Keyan Xu, Qingfu Zhu, Wanxiang Che, and Yang Deng. Learning-to-context slope: Evaluating in-context learning effectiveness beyond performance illusions, 2025.
- [53] Futing Wang, Jianhao Yan, Yue Zhang, and Tao Lin. ELICIT: LLM augmentation via external in-context capability. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [54] Lean Wang, Lei Li, Damai Dai, Deli Chen, Hao Zhou, Fandong Meng, Jie Zhou, and Xu Sun. Label words are anchors: An information flow perspective for understanding in-context learning. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9840–9855, Singapore, December 2023. Association for Computational Linguistics.
- [55] Shumin Wang, Yuexiang Xie, Bolin Ding, Jinyang Gao, and Yanyong Zhang. Language adaptation of large language models: An empirical study on LLaMA2. In Owen Rambow, Leo Wanner, Marianna Apidianaki, Hend Al-Khalifa, Barbara Di Eugenio, and Steven Schockaert, editors, *Proceedings of the 31st International Conference on Computational Linguistics*, pages 7195–7208, Abu Dhabi, UAE, January 2025. Association for Computational Linguistics.
- [56] Song Wang, Zihan Chen, Chengshuai Shi, Cong Shen, and Jundong Li. Mixture of demonstrations for in-context learning. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [57] Xubin Wang, Jianfei Wu, Yuan Yichen, Deyu Cai, Mingzhe Li, and Weijia Jia. Demonstration selection for in-context learning via reinforcement learning. In *Forty-second International Conference on Machine Learning*, 2025.
- [58] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. MMLU-pro: A more robust and challenging multi-task language understanding benchmark. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.
- [59] Jerry Wei, Jason Wei, Yi Tay, Dustin Tran, Albert Webson, Yifeng Lu, Xinyun Chen, Hanxiao Liu, Da Huang, Denny Zhou, and Tengyu Ma. Larger language models do in-context learning differently, 2024.
- [60] Kevin Christian Wibisono and Yixin Wang. From unstructured data to in-context learning: Exploring what tasks can be learned and when. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. OpenReview (NeurIPS 2024).

- [61] Noam Wies, Yoav Levine, and Amnon Shashua. The learnability of in-context learning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [62] Yuan Wu, Diana Inkpen, and Ahmed El-Roby. Maximum batch frobenius norm for multi-domain text classification. In *IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2022, Virtual and Singapore, 23-27 May 2022*, pages 3763–3767. IEEE, 2022.
- [63] Vikas Yadav, Steven Bethard, and Mihai Surdeanu. Quick and (not so) dirty: Unsupervised selection of justification sentences for multi-hop question answering. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2578–2589, Hong Kong, China, November 2019. Association for Computational Linguistics.
- [64] Tong Yang, Yu Huang, Yingbin Liang, and Yuejie Chi. In-context learning with representations: Contextual generalization of trained transformers. In *ICML 2024 Workshop on Theoretical Foundations of Foundation Models*, 2024.
- [65] Zhao Yang, Yuanzhe Zhang, Dianbo Sui, Cao Liu, Jun Zhao, and Kang Liu. Representative demonstration selection for in-context learning with two-stage determinantal point process. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5443–5456, Singapore, December 2023. Association for Computational Linguistics.
- [66] Xi Ye, Srinivasan Iyer, Asli Celikyilmaz, Veselin Stoyanov, Greg Durrett, and Ramakanth Pasunuru. Complementary explanations for effective in-context learning. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Findings of the Association for Computational Linguistics: ACL 2023*, pages 4469–4484, Toronto, Canada, July 2023. Association for Computational Linguistics.
- [67] Ruiqi Zhang, Spencer Frei, and Peter L. Bartlett. Trained transformers learn linear models in-context. *Journal of Machine Learning Research*, 25(49):1–55, 2024.
- [68] Xingxuan Zhang, Haoran Wang, Jiansheng Li, Yuan Xue, Shikai Guan, Renzhe Xu, Hao Zou, Han Yu, and Peng Cui. Understanding the generalization of in-context learning in transformers: An empirical study. In *International Conference on Learning Representations (ICLR)*, 2025. OpenReview (ICLR 2025).
- [69] Yedi Zhang, Aaditya K Singh, Peter E. Latham, and Andrew M Saxe. Training dynamics of in-context learning in linear attention. In *Forty-second International Conference on Machine Learning*, 2025.
- [70] Siyan Zhao, Tung Nguyen, and Aditya Grover. Probing the decision boundaries of in-context learning in large language models. In *NeurIPS (Poster)*, 2024. OpenReview (NeurIPS 2024).
- [71] Yuxiang Zhou, Jiazheng Li, Yanzheng Xiang, Hanqi Yan, Lin Gui, and Yulan He. The mystery of in-context learning: A comprehensive survey on interpretation and analysis. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 14365–14378, Miami, Florida, USA, November 2024. Association for Computational Linguistics.

## A Appendix

### A.1 Proof

#### A.1.1 Proof of Equation 2

*Proof.* Based on Equation 1 and [67], the predicted answer of LSA is:

$$\hat{y}_{query} = ((w_{21}^{PV})^\top \quad w_{22}^{PV}) \cdot \left( \frac{EE^\top}{N} \right) \cdot \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x,$$

where  $N$  is the number of demonstrations. For  $N = 1$ , the matrix product  $EE^\top$  can be expressed compactly as  $EE^\top = dd^\top + qq^\top$ . Substituting this into the equation gives:

$$\hat{y}_{query} = ((w_{21}^{PV})^\top \quad w_{22}^{PV})(dd^\top + qq^\top) \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x$$

We can expand this expression and separate the terms that depend on  $d$ :

$$\begin{aligned} \hat{y}_{query} &= \underbrace{((w_{21}^{PV})^\top \quad w_{22}^{PV})(dd^\top) \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x}_{\text{Term depending on } d} + \\ &\quad \underbrace{((w_{21}^{PV})^\top \quad w_{22}^{PV})(qq^\top) \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x}_{\text{Constant w.r.t. } d} \end{aligned}$$

The term depending on  $d$  can be rewritten using the property of scalar products:

$$\begin{aligned} &((w_{21}^{PV})^\top \quad w_{22}^{PV})dd^\top \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x \\ &= (((w_{21}^{PV})^\top \quad w_{22}^{PV})d) \left( d^\top \begin{pmatrix} W_{11}^{KQ} \\ (w_{21}^{KQ})^\top \end{pmatrix} q_x \right) \end{aligned}$$

This is a quadratic form of the vector  $d$ . To compute its gradient, we use the vector calculus identity  $\nabla_z((a^\top z)(b^\top z)) = a(b^\top z) + b(a^\top z)$ . We set:

$$\begin{aligned} a &= \begin{pmatrix} w_{21}^{PV} \\ w_{22}^{PV} \end{pmatrix} \\ b &= \begin{pmatrix} W_{11}^{KQ} q_x \\ (w_{21}^{KQ})^\top q_x \end{pmatrix} \end{aligned}$$

The second term in the expression for  $\hat{y}_{query}$  is constant with respect to  $d$ , so its gradient is zero. Applying the identity to the first term yields the gradient of  $\hat{y}_{query}$  with respect to  $d$ :

$$\begin{aligned} \nabla_d \hat{y}_{query} &= \begin{pmatrix} w_{21}^{PV} \\ w_{22}^{PV} \end{pmatrix} \left( \begin{pmatrix} d_x \\ d_y \end{pmatrix}^\top \begin{pmatrix} W_{11}^{KQ} q_x \\ (w_{21}^{KQ})^\top q_x \end{pmatrix} \right) + \\ &\quad \begin{pmatrix} W_{11}^{KQ} q_x \\ (w_{21}^{KQ})^\top q_x \end{pmatrix} \left( \begin{pmatrix} w_{21}^{PV} \\ w_{22}^{PV} \end{pmatrix}^\top \begin{pmatrix} d_x \\ d_y \end{pmatrix} \right) \\ &= \begin{pmatrix} w_{21}^{PV} \\ w_{22}^{PV} \end{pmatrix} \left( d_x^\top W_{11}^{KQ} q_x + d_y (w_{21}^{KQ})^\top q_x \right) + \\ &\quad \begin{pmatrix} W_{11}^{KQ} q_x \\ (w_{21}^{KQ})^\top q_x \end{pmatrix} \left( (w_{21}^{PV})^\top d_x + d_y w_{22}^{PV} \right) \end{aligned}$$

This completes the derivation.  $\square$

### A.2 Proof of Lemma 1

*Proof.* We prove by induction on the layer index  $l$ .

**Base case ( $l = 0$ ).** The two inequalities in the statement are satisfied by assumption.

**Induction step.** Assume the inequalities hold for some layer  $l - 1$  ( $1 \leq l \leq L$ ), i.e.

$$\begin{aligned} \|W^{PV,(l-1)}d_1^{(l-1)}\| &\geq \|W^{PV,(l-1)}d_2^{(l-1)}\| \\ \|(d_1^{(l-1)})^\top W^{KQ,(l-1)}q^{(l-1)}\| &\geq \\ \|(d_2^{(l-1)})^\top W^{KQ,(l-1)}q^{(l-1)}\|. \end{aligned}$$

Apply (5) to both demonstrations and use the strict monotonicity of  $g_l$ :

$$\begin{aligned} &\|W^{PV,(l)}d_1^{(l)}\| \\ &= g_l(\|W^{PV,(l-1)}d_1^{(l-1)}\|) \\ &\geq g_l(\|W^{PV,(l-1)}d_2^{(l-1)}\|) \\ &= \|W^{PV,(l)}d_2^{(l)}\|. \end{aligned}$$

Similarly, apply (6) and the strict monotonicity of  $h_l$ :

$$\begin{aligned} &\|(d_1^{(l)})^\top W^{KQ,(l)}q^{(l)}\| \\ &= h_l(\|(d_1^{(l-1)})^\top W^{KQ,(l-1)}q^{(l-1)}\|) \\ &\geq h_l(\|(d_2^{(l-1)})^\top W^{KQ,(l-1)}q^{(l-1)}\|) \\ &= \|(d_2^{(l)})^\top W^{KQ,(l)}q^{(l)}\|. \end{aligned}$$

Thus the claim holds for layer  $l$ . By induction it holds for all layers  $l = 0, \dots, L$ .  $\square$

### A.2.1 Proof of Theorem 1

*Proof.* Based on Equation 4, we can derive:

$$\begin{aligned} &\frac{\partial \hat{q}_y^{(l_1)}}{\partial d^{(0)}} / \frac{\partial \hat{q}_y^{(l_2)}}{\partial d^{(0)}} \\ &= \left( \frac{\partial \hat{q}_y^{(l_1)}}{\partial E^{(l_1-1)}} \times \frac{\partial E^{(1)}}{\partial d^{(0)}} \times \prod_{i=2}^{l_1} \frac{\partial E^{(i)}}{\partial E^{(i-1)}} \right) / \\ &\quad \left( \frac{\partial \hat{q}_y^{(l_2)}}{\partial E^{(l_2-1)}} \times \frac{\partial E^{(1)}}{\partial d^{(0)}} \times \prod_{i=2}^{l_2} \frac{\partial E^{(i)}}{\partial E^{(i-1)}} \right) \\ &= \left( \frac{\partial \hat{q}_y^{(l_1)}}{\partial E^{(l_1-1)}} / \frac{\partial \hat{q}_y^{(l_2)}}{\partial E^{(l_2-1)}} \right) \times \prod_{i=l_2+1}^{l_1} \frac{\partial E^{(i)}}{\partial E^{(i-1)}} \end{aligned}$$

Consider the total derivative of  $f_{\text{LSA}}$ :

$$\begin{aligned} df_{\text{LSA}} &= dE + W^{PV} dE E^\top W^{KQ} E + \\ &\quad W^{PV} E (E^\top W^{KQ} dE + dE^\top W^{KQ} E) \end{aligned}$$

It can be observed that the coefficients of the differential terms are consistent with those in Definition 1. Therefore, from  $d_1^{(l)} \succcurlyeq_{q^{(l)}; \theta^{(l)}} d_2^{(l)}$ , we know that:

$$\frac{\partial E^{(l)}}{\partial E^{(l-1)}}(E_1; \theta) \geq \frac{\partial E^{(l)}}{\partial E^{(l-1)}}(E_2; \theta), \forall l \in \{l_2 + 1, \dots, l_1\}$$

Therefore, we can conclude that:

$$\frac{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_1)}(E_1; \theta)\|}{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_1)}(E_2; \theta)\|} \geq \frac{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_2)}(E_1; \theta)\|}{\|\nabla_{d^{(0)}} \hat{q}_y^{(l_2)}(E_2; \theta)\|}$$

$\square$

---

**Prompt of Inference**

---

{task}

Below are some examples

---

—  
{demo}

---

—  
Based on the above instruction and examples, solve the following problem.  
{question}

---

Table 4: The prompt of the inference.

### A.3 Prompt

In this section, we present the inference prompt of our main experiment, as shown in Table 4. The task definition we used is same to the previous works [11, 1].

### A.4 Dataset

Dataset	Test Set	Demonstration
GSM8K	1319	7473
MATH	500	7496
ARC-Challenge	1172	1119
MMLU-Pro	1000	70
Amazon Review	200	1800

Table 5: The scales of test set and demonstrations of each dataset.

In this section, we detail the datasets used in our study. Table 5 summarizes the scale of the test set and demonstrations for each.

**GSM8K** GSM8K [10] is a high-quality collection of elementary school-level math problems. We utilize its training set directly as the demonstration pool.

**MATH** The MATH dataset [17] consists of challenging high school competition-level math problems in fields like algebra, probability, and geometry. Following the approach of [25], we evaluate GRADS on a random sample of 500 problems. The demonstrations are drawn from the official training set.

**ARC-Challenge** The ARC-Challenge [63] is a question-answering dataset with difficult, science-focused questions. For this dataset, the training set is used as our demonstration pool.

**MMLU-Pro** MMLU-Pro [58] serves as a multi-task benchmark for the comprehensive evaluation of LLMs on professional domain knowledge and complex reasoning. As the dataset is only divided into validation and test sets, we use the validation set as our demonstration pool and conduct evaluations on the test set.

**Amazon Review** The Amazon Review dataset [34], containing a vast amount of user ratings and reviews, is widely used for research in sentiment analysis and recommender systems. Due to the immense size, we select the *Health and Personal Care* category for testing, while using the *All Beauty*, *Digital Music*, and *Software* categories to form the demonstration pool.

## A.5 Baseline

**BM25** BM25 is a classic sparse retrieval method based on the probabilistic relevance framework, serving as an extension of TF-IDF. In the context of ICL, it treats the test query as a search query and the candidate demonstrations as documents. It ranks demonstrations by scoring them based on the frequency and distribution of query terms, without considering word order or deep semantics. The top-K scored demonstrations are then selected as the in-context demonstrations. This “bag-of-words” approach is known for its computational efficiency and serves as a strong baseline.

**Cosine Similarity** Cosine Similarity is a popular and effective strategy for ICL demonstration selection that aims to find demonstrations semantically similar to the test query. This method involves encoding the test query and candidate demonstrations into high-dimensional vectors (embeddings) using a pre-trained sentence encoder, such as BERT. The semantic relevance between the query and each demonstration is then measured by computing the cosine similarity of their respective vectors. The demonstrations with the highest similarity scores are selected to form the context. This dense retrieval approach generally captures semantic nuances better than sparse methods like BM25, but its performance depends on the quality of the underlying embedding model.

**MMR** Maximal Marginal Relevance (MMR) is a selection strategy designed to balance the relevance of demonstrations to the query with the diversity within the selected set. The rationale is that selecting only the nearest neighbors (most relevant examples) can result in a set of overly similar demonstrations, which may limit the variety of reasoning processes shown to the model. MMR addresses this by iteratively selecting demonstrations that maximize a combined score of relevance to the query and dissimilarity from the examples already chosen. This approach aims to create a set of demonstrations that are not only relevant but also complementary, using diversity as a proxy for complementarity to improve ICL performance.

**MoD** Mixture of Demonstrations (MoD) is a framework designed to overcome the challenges of a large search space and suboptimal retriever optimization in ICL demonstration selection. The core idea is to partition the entire demonstration pool into distinct groups, typically using K-means clustering on sentence embeddings. Each group is governed by a dedicated “expert”, a unique retriever model trained specifically for that partition. During inference, these experts collaboratively retrieve demonstrations for a given query, with the final set being an aggregation of examples selected by the most relevant experts. This “mixture of experts” approach reduces the search complexity while ensuring the selected demonstrations are diverse and effective.

## A.6 Additional Experiment Results

### A.6.1 Scale of Effective and Ineffective Data

The number of effective and ineffective demonstrations we sampled is shown in Table 6.

### A.6.2 Detailed Gradient Flow under Each Setting

In this part, we present the gradient flow under each setting from Figure 4 to Figure 23.

Model	Dataset	Effective	Ineffective
<b>Llama2-7b</b>	GSM8K	225	1030
	MATH	18	452
	ARC-Challenge	355	674
	MMLU-Pro	108	843
	Amazon	38	134
<b>Llama3.1-8b</b>	GSM8K	84	210
	MATH	52	280
	ARC-Challenge	86	212
	MMLU-Pro	174	466
	Amazon	26	63
<b>Llama-R1-8b</b>	GSM8K	41	519
	MATH	27	146
	ARC-Challenge	77	198
	MMLU-Pro	203	550
	Amazon	20	102
<b>Qwen3-8b</b>	GSM8K	33	99
	MATH	31	111
	ARC-Challenge	28	110
	MMLU-Pro	251	368
	Amazon	23	77

Table 6: The number of effective and ineffective demonstrations under each setting.

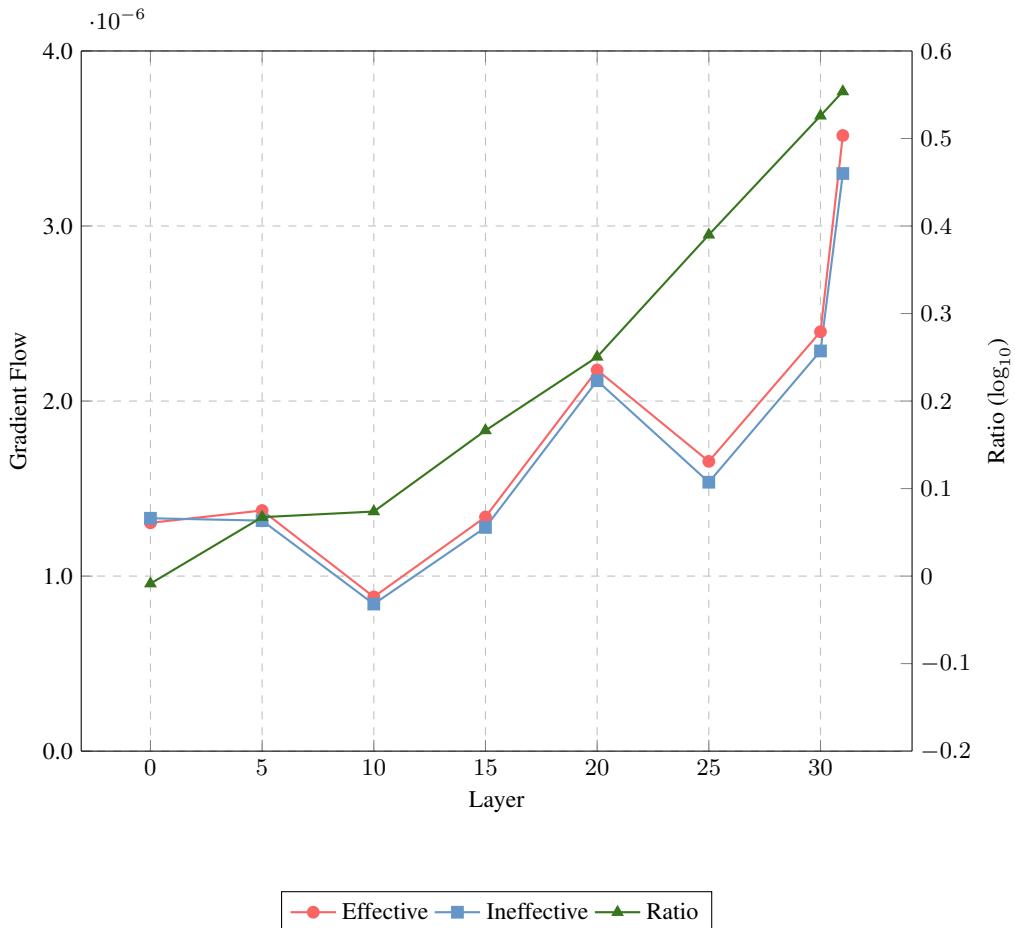


Figure 4: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama2-7b on GSM8K.

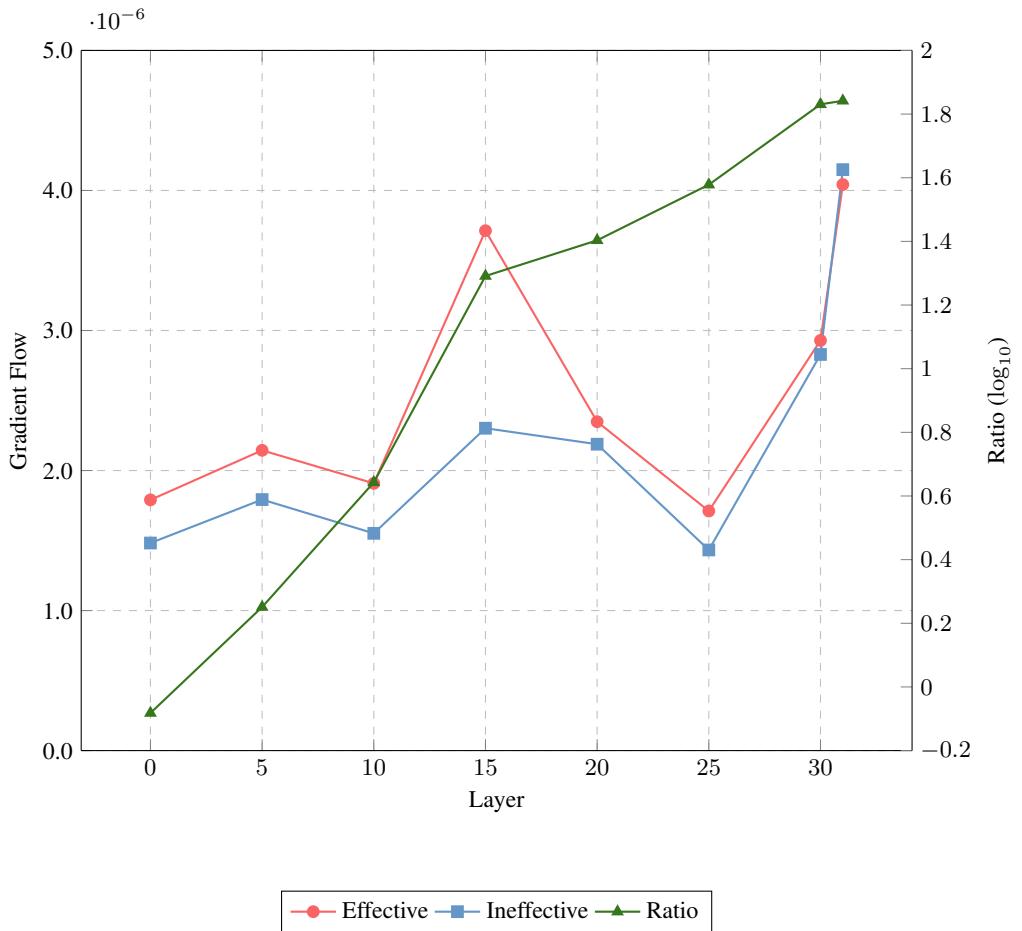


Figure 5: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama2-7b on MATH.

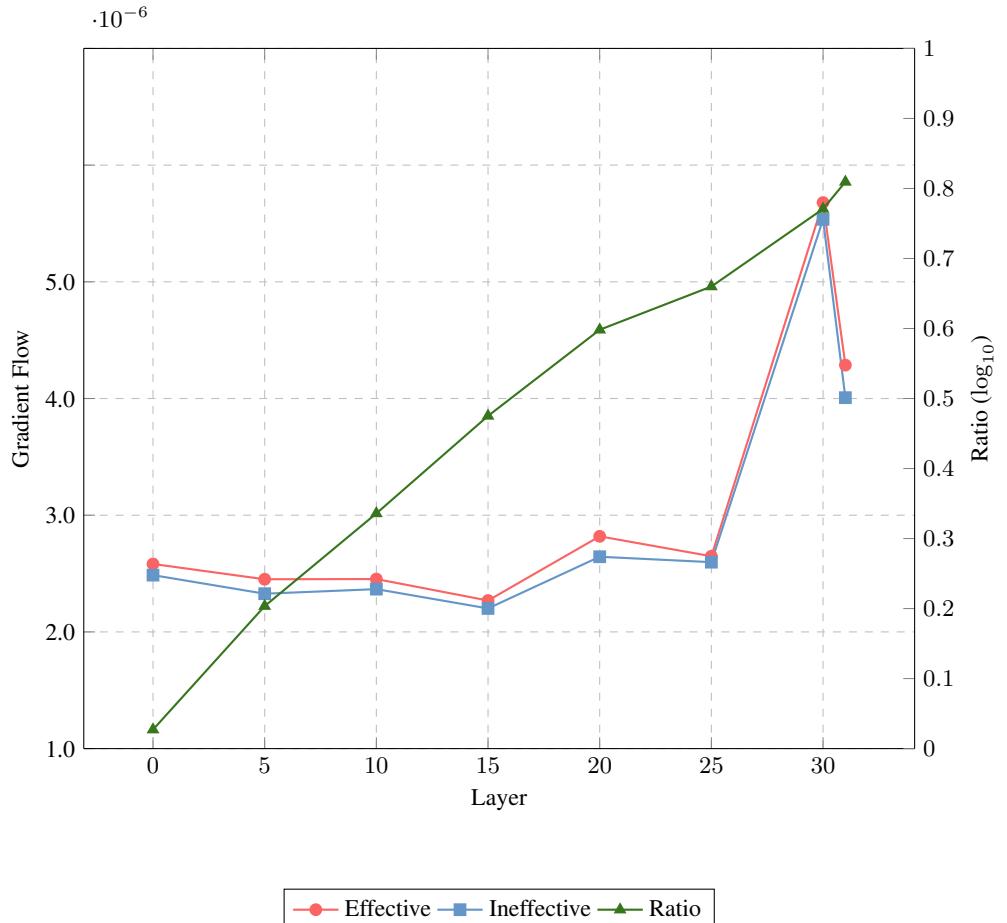


Figure 6: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama2-7b on ARC-Challenge.

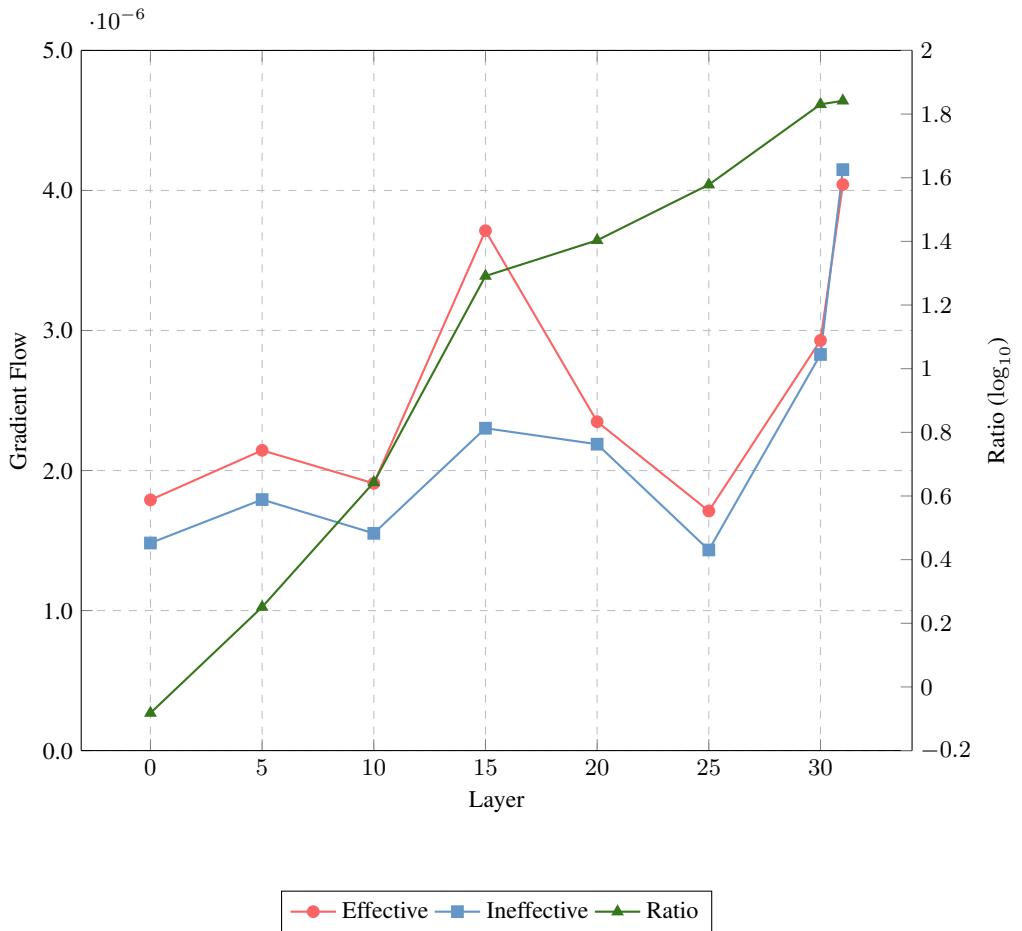


Figure 7: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama2-7b on MMLU-Pro.

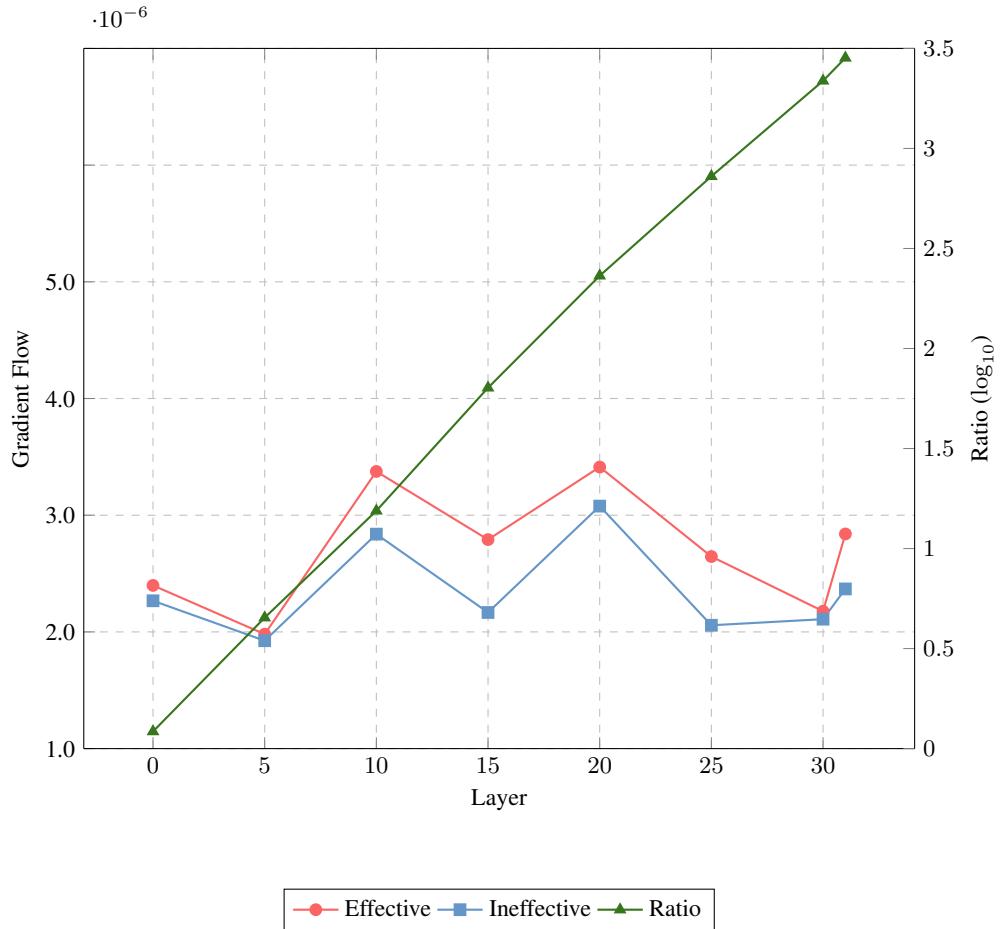


Figure 8: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama2-7b on Amazon Review.

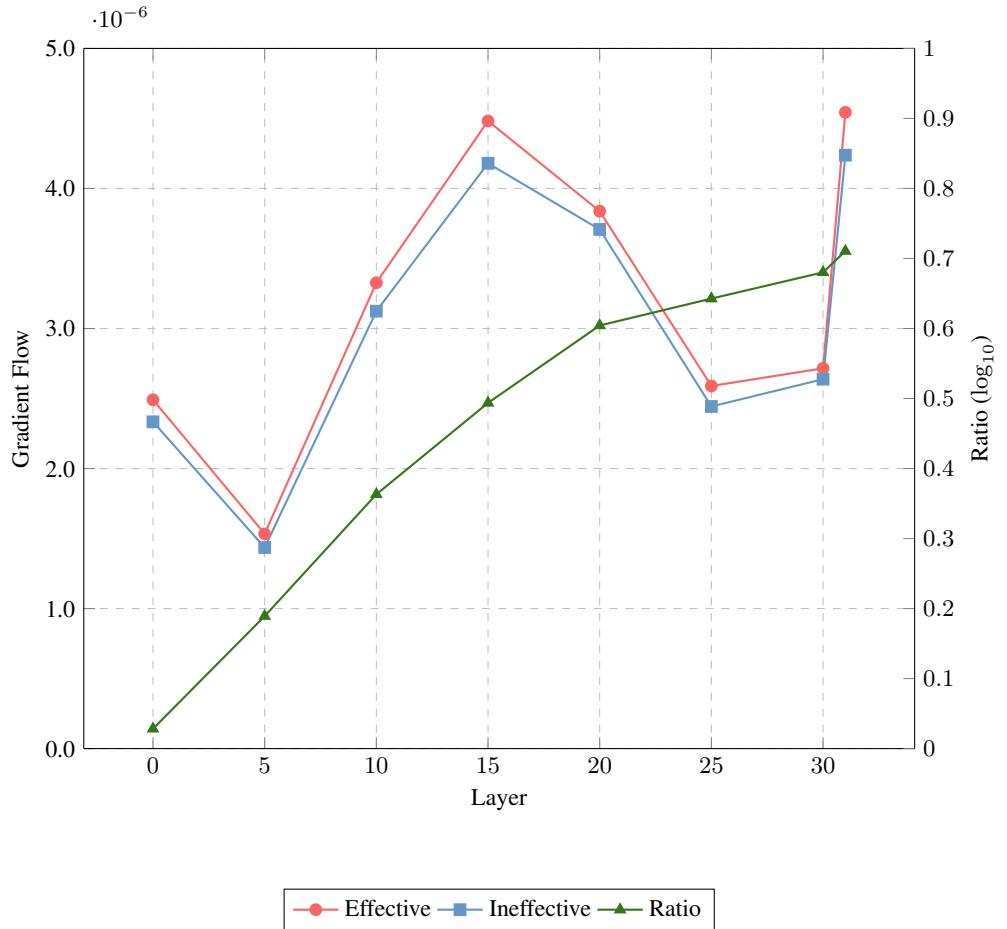


Figure 9: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama3.1-8b on GSM8K.

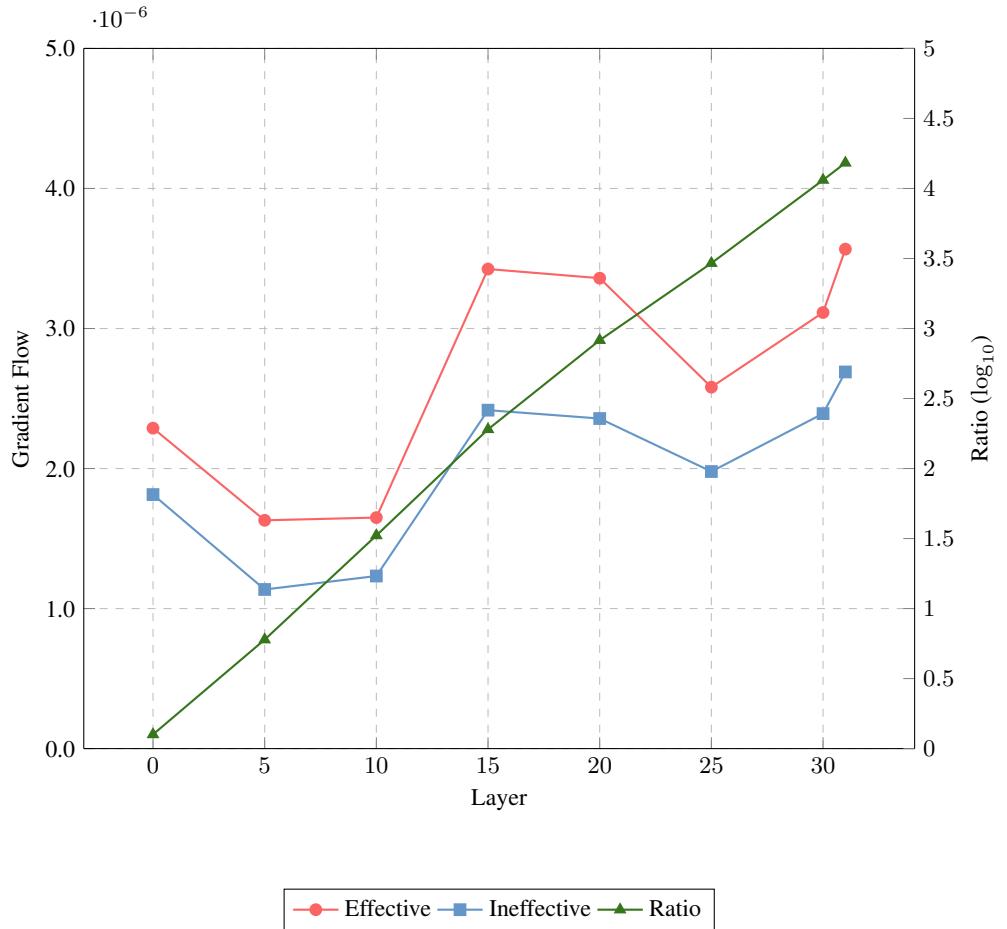


Figure 10: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama3.1-8b on MATH.

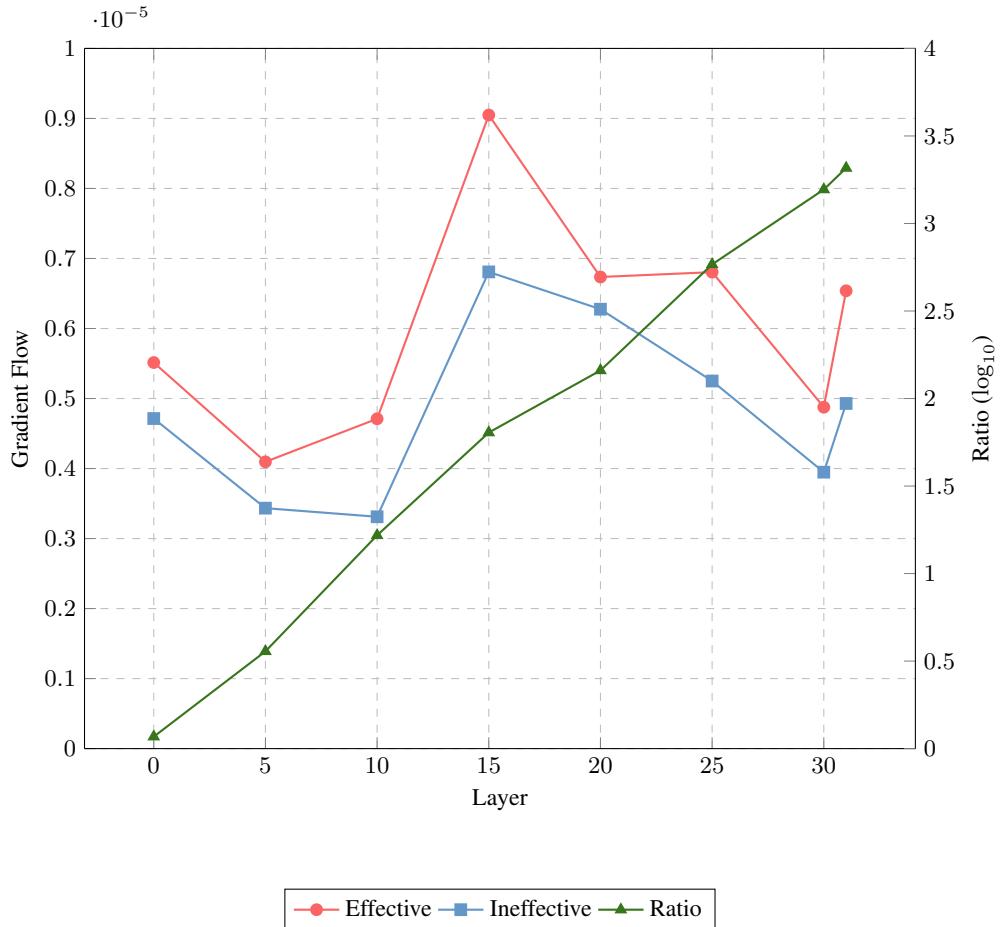


Figure 11: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama3.1-8b on ARC-Challenge.

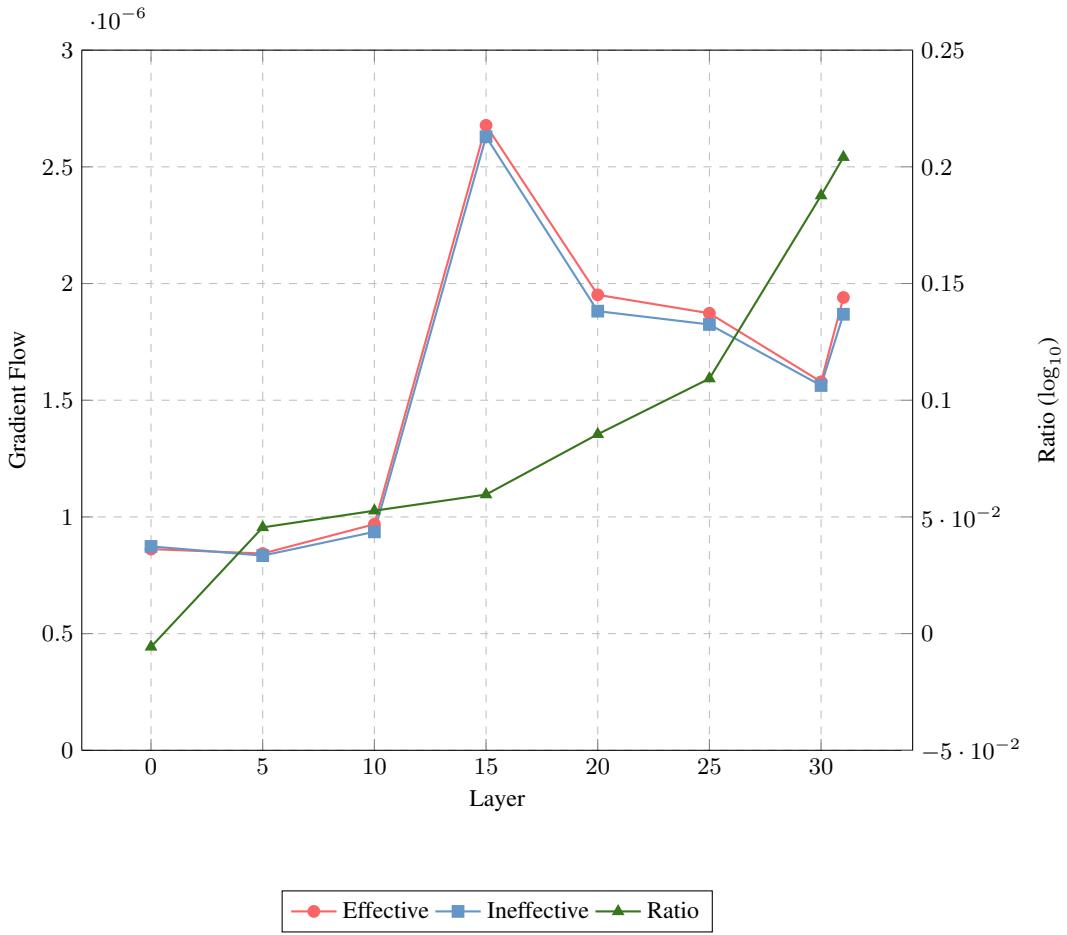


Figure 12: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama3.1-8b on MMLU-Pro.

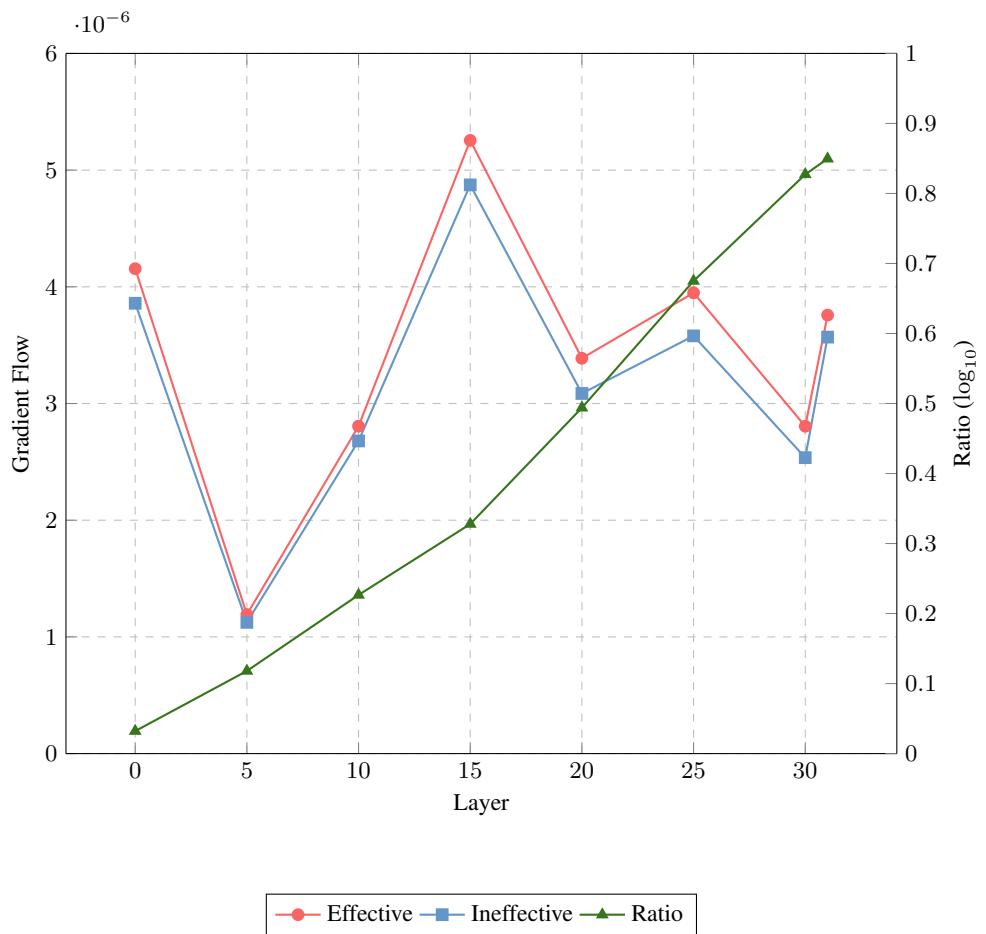


Figure 13: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama3.1-8b on Amazon Review.

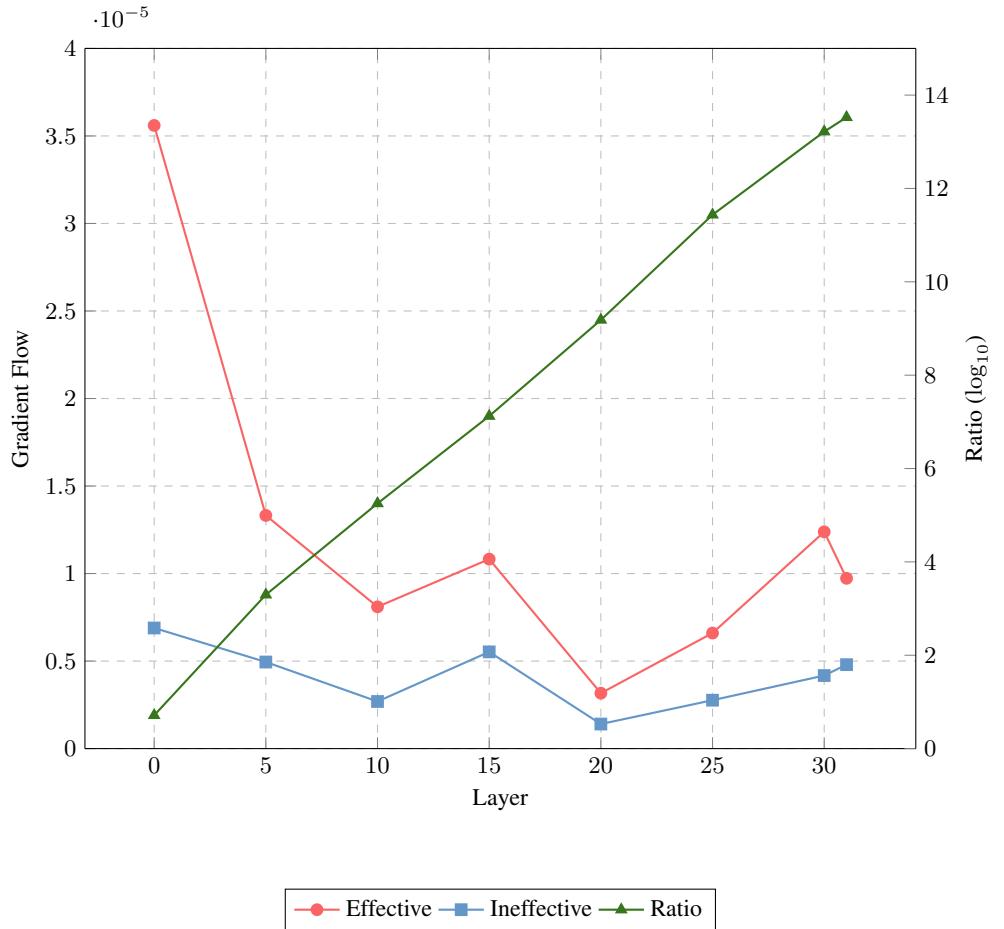


Figure 14: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama-R1-8b on GSM8K.

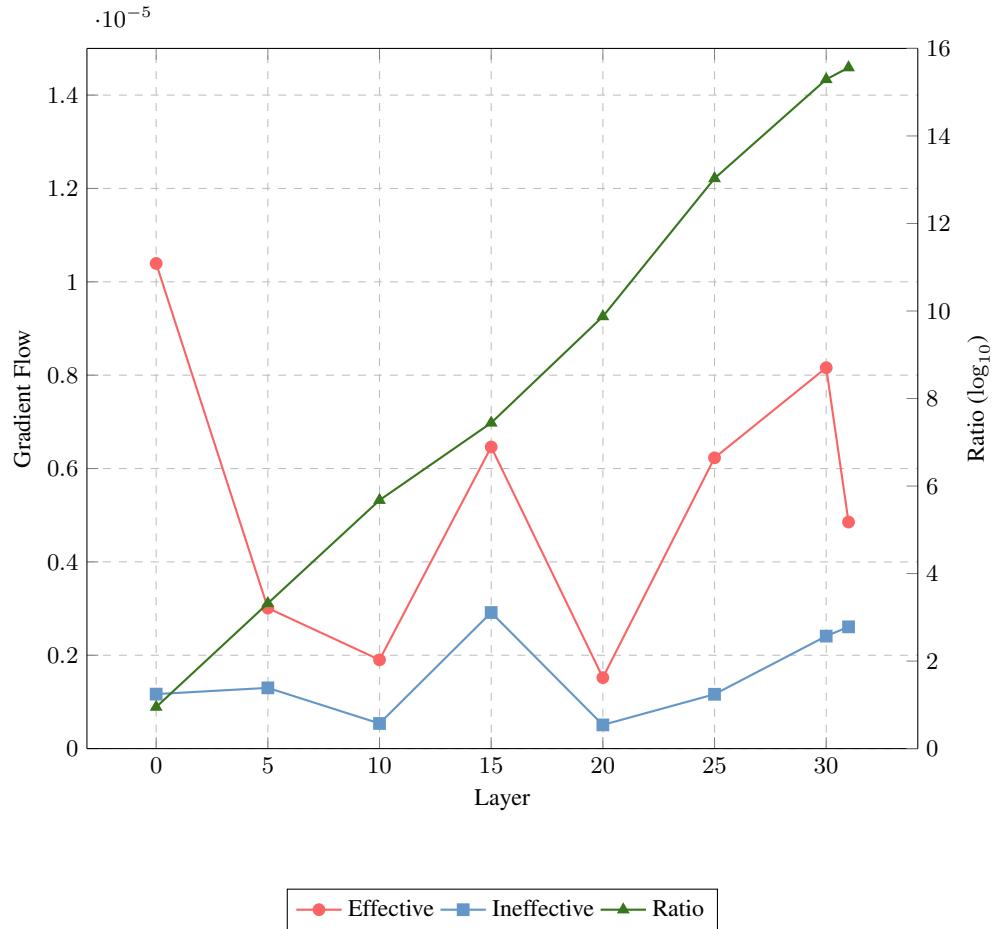


Figure 15: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama-R1-8b on MATH.

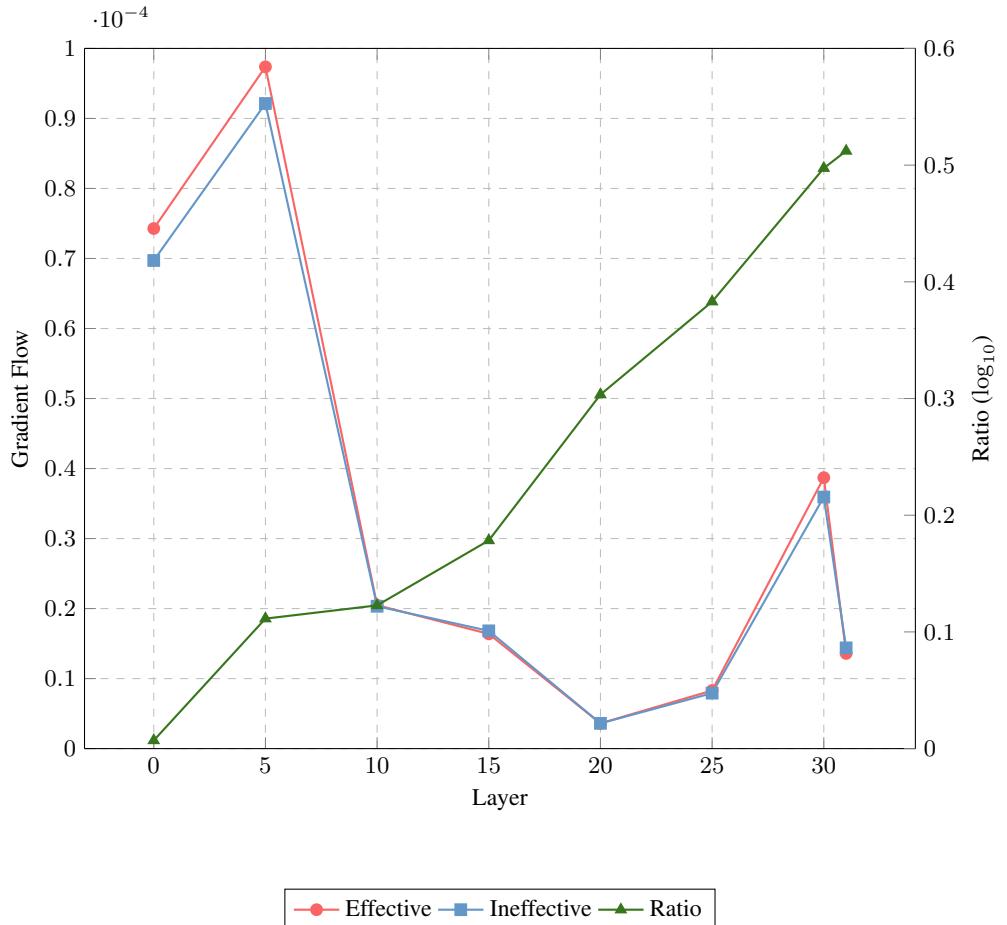


Figure 16: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama-R1-8b on ARC-Challenge.

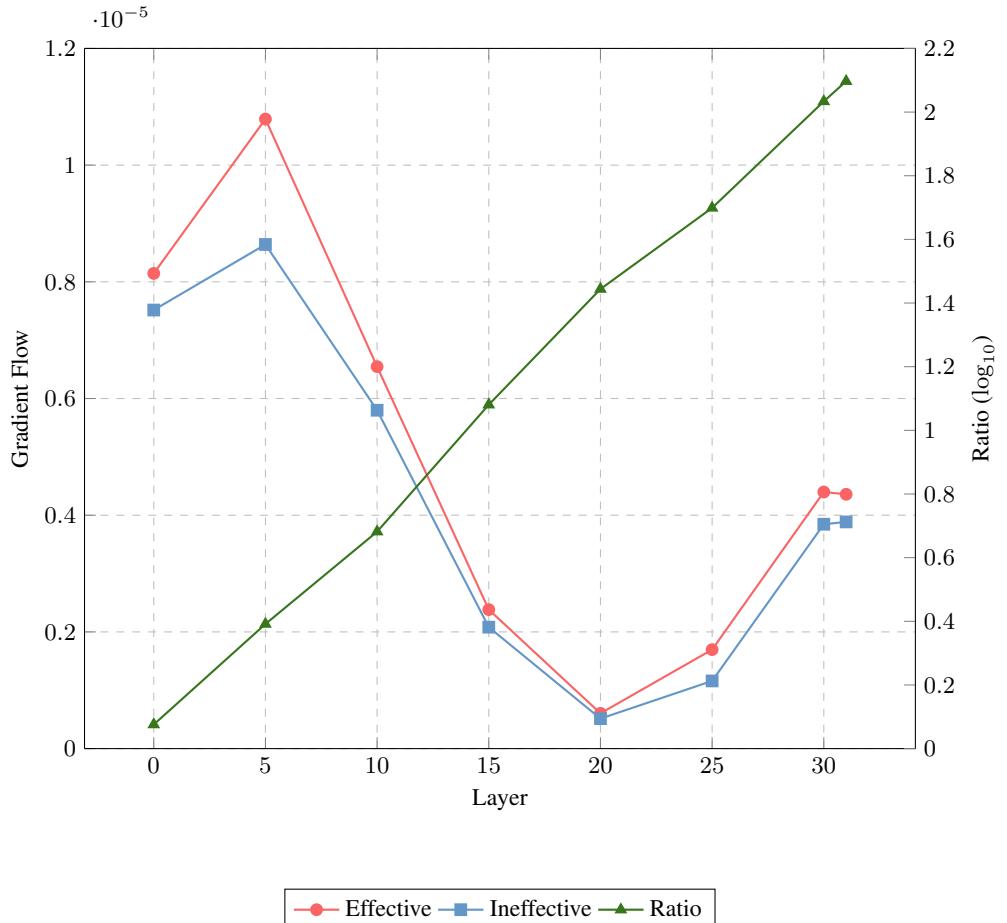


Figure 17: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama-R1-8b on MMLU-Pro.

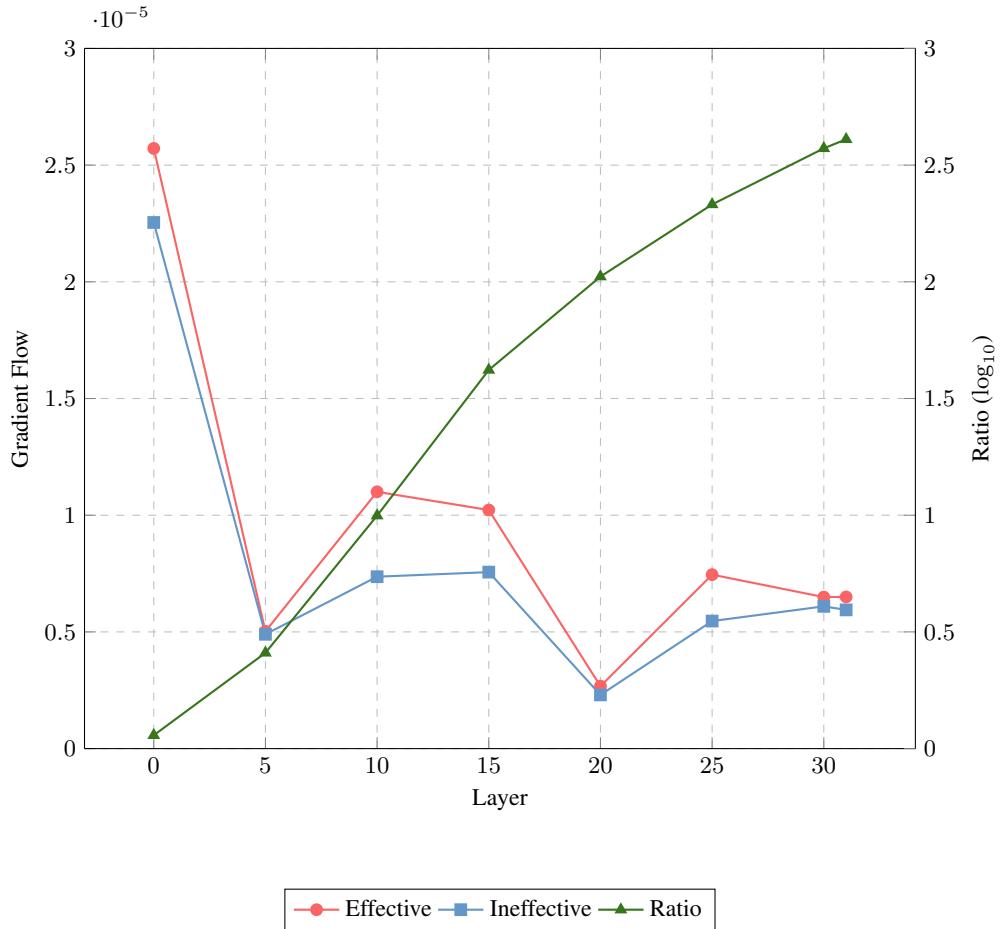


Figure 18: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Llama-R1-8b on Amazon Review.

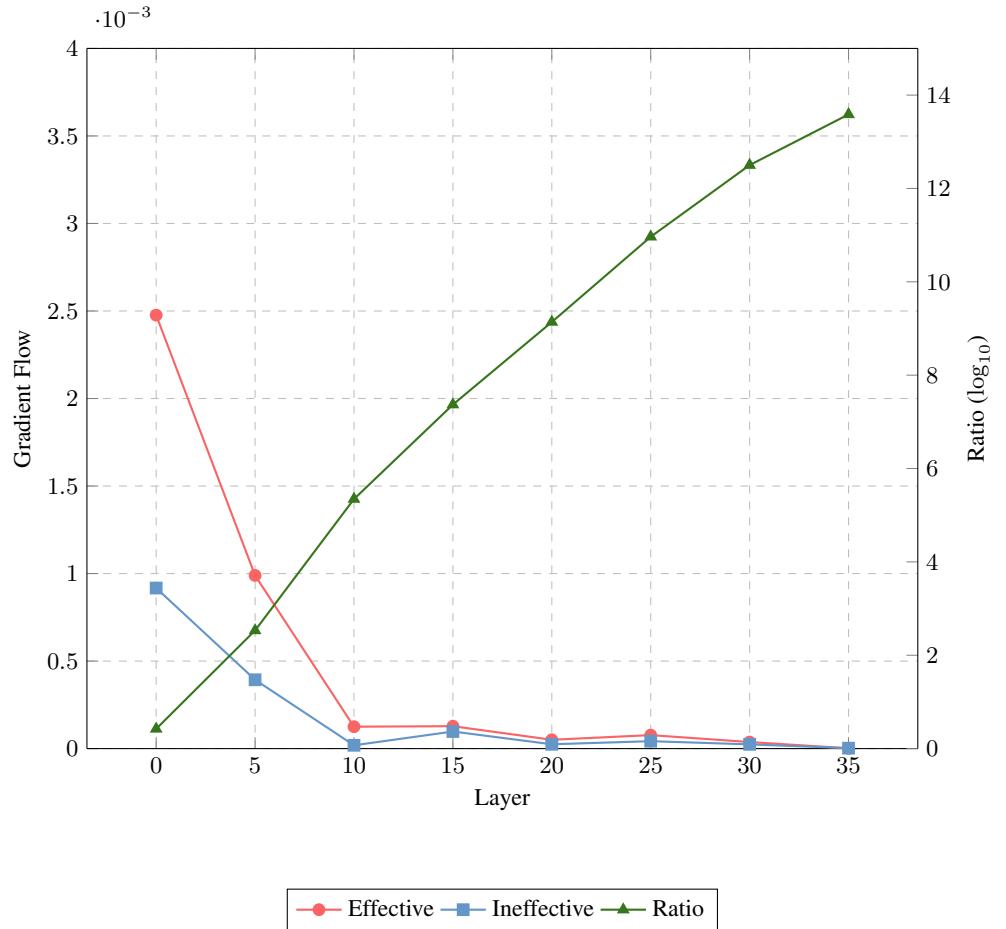


Figure 19: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Qwen3-8b on GSM8K.

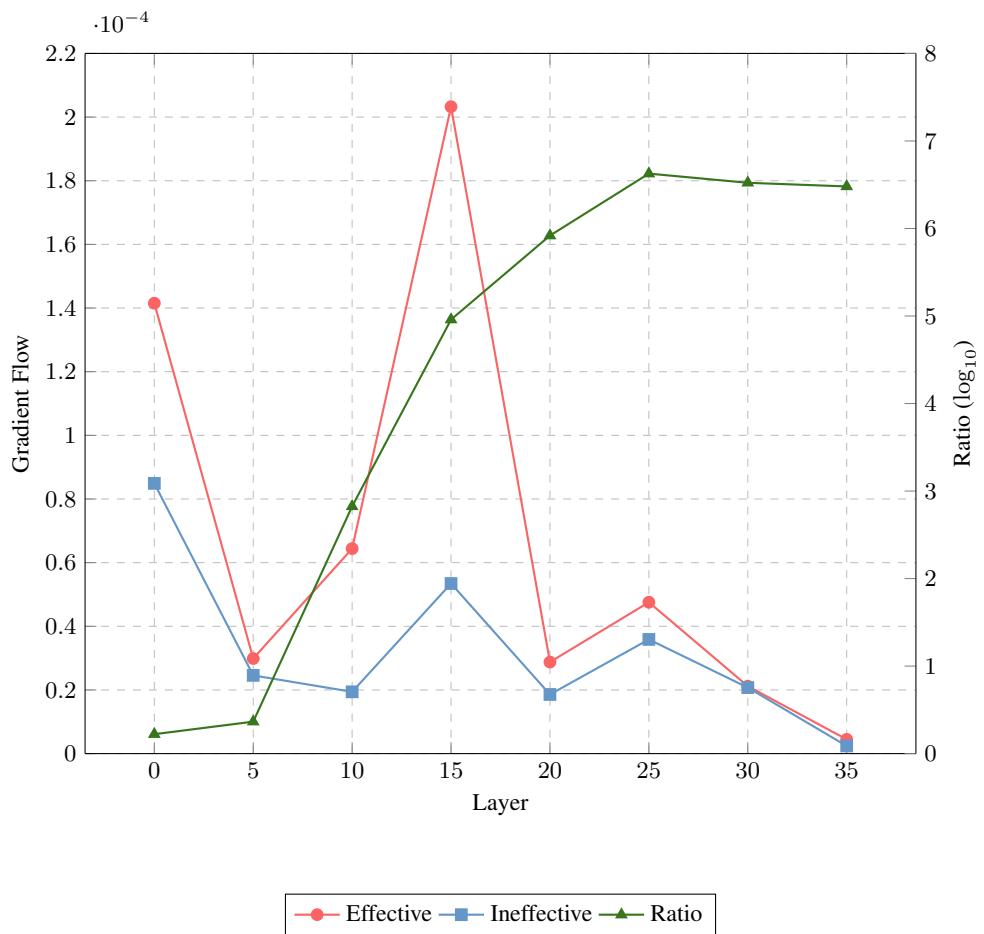


Figure 20: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Qwen3-8b on MATH.

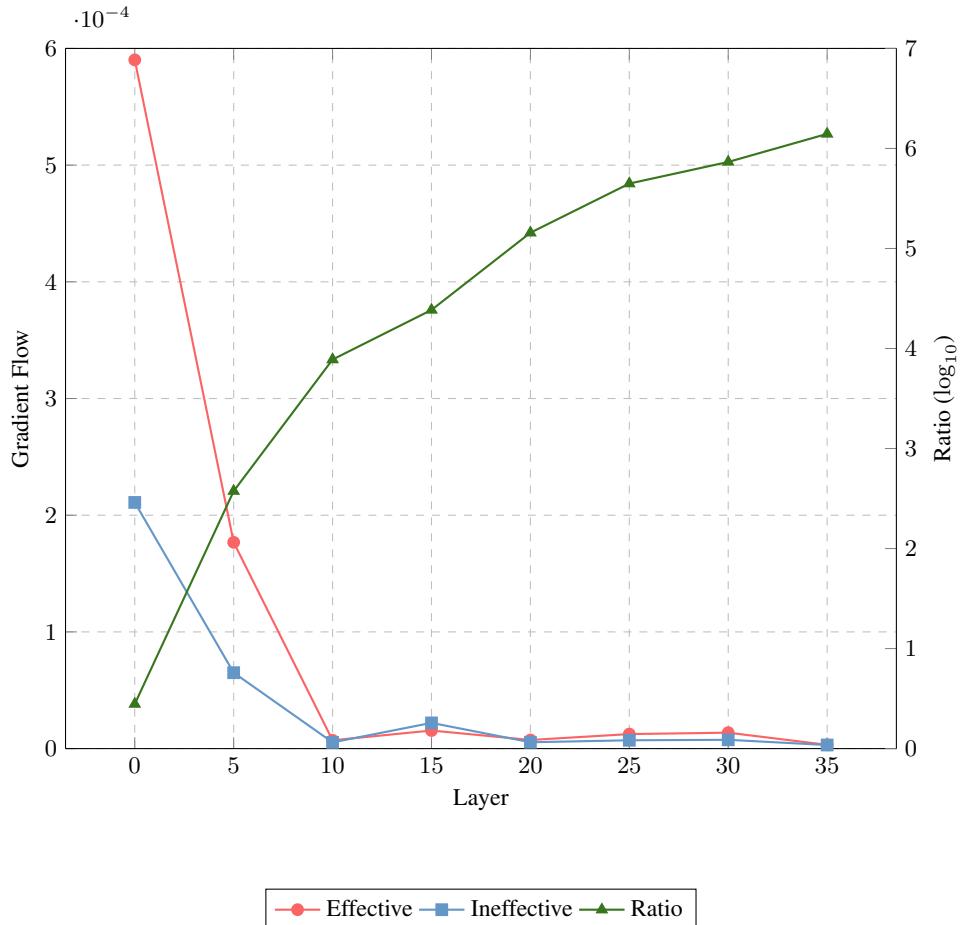


Figure 21: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Qwen3-8b on ARC-Challenge.

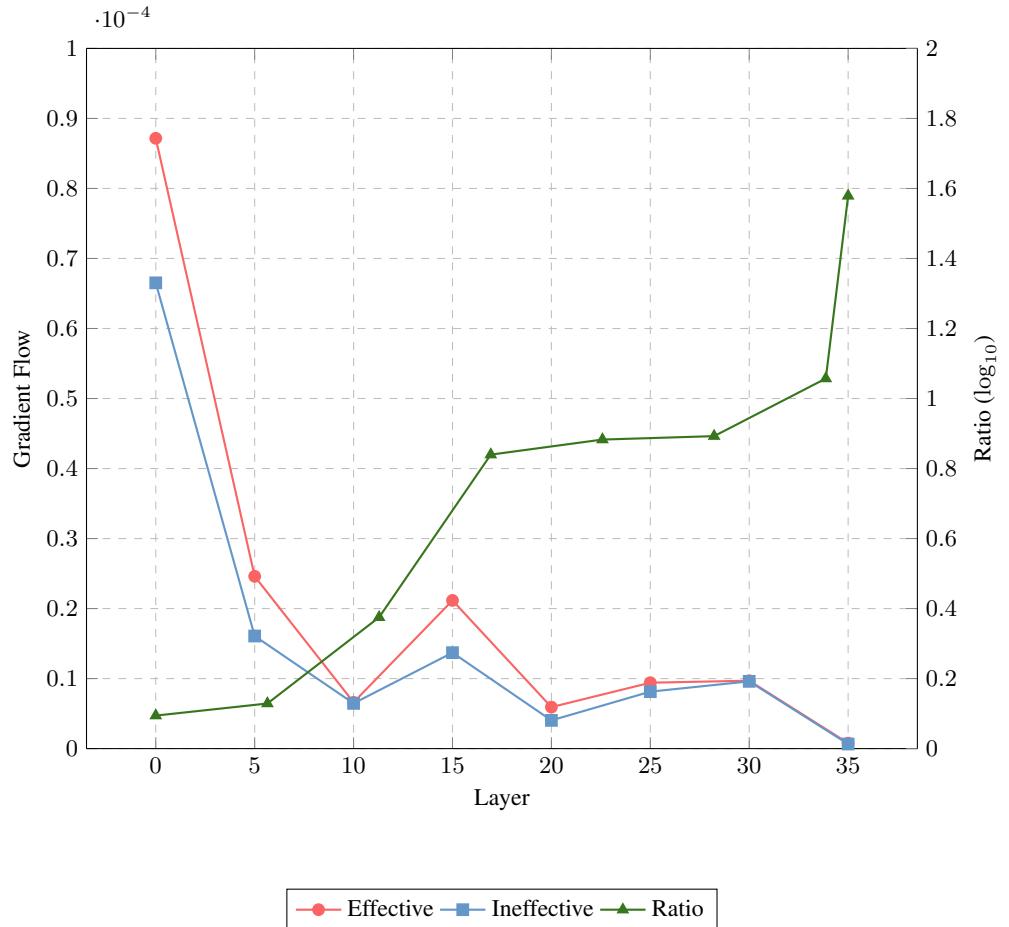


Figure 22: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Qwen3-8b on MMLU-Pro.

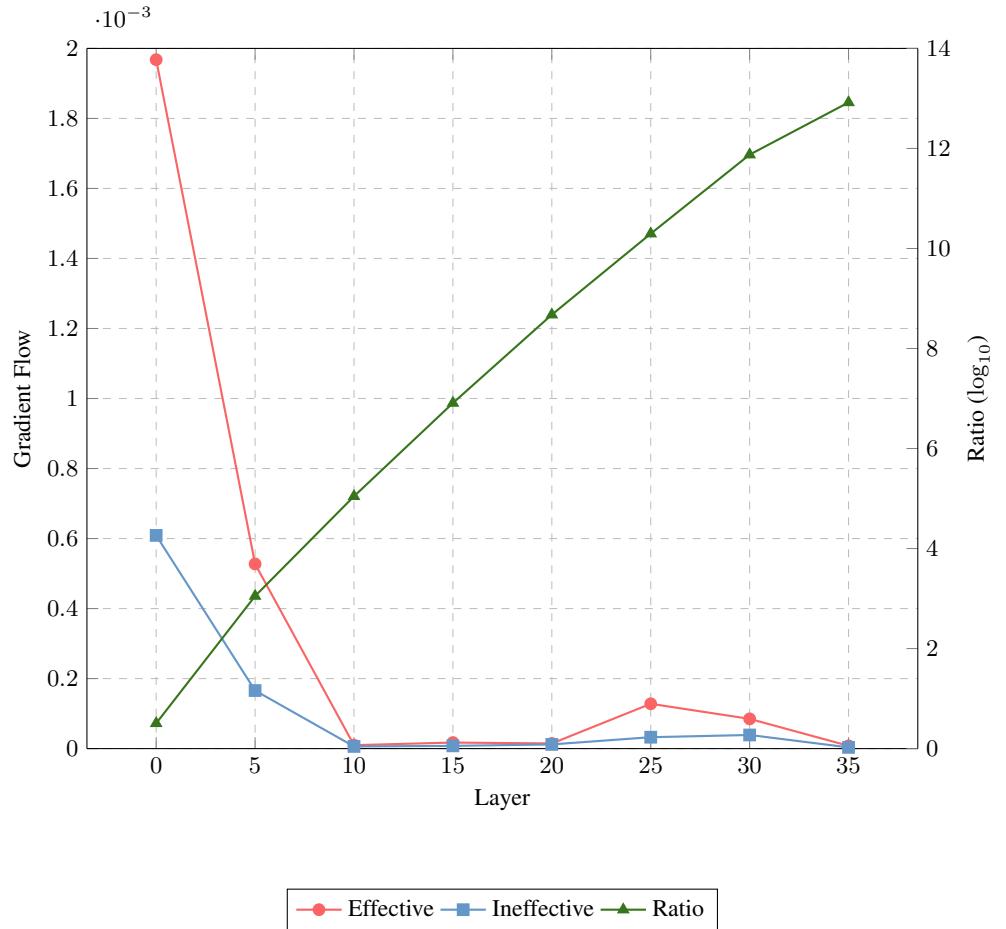


Figure 23: The average gradient flow (left y-axis) and the ratio between the effective and the ineffective (right y-axis) under each layer of Qwen3-8b on Amazon Review.